

VeriExploit: Automatic Bug Reproduction in Smart Contracts via LLMs and Formal Methods

Chenfeng Wei
Chenfeng.wei@manchester.ac.uk
The University of Manchester

Co-Authors




ESBMC



Mr.
Chenfeng Wei



Mr.
Shiyu Cai



Mr.
Yiannis Charalambous



Mr.
Tong Wu



Dr.
Sangharatna Godbole



Dr.
Lucas C. Cordeiro

Motivation

Reproducing found bugs can be hard

- Bug finders (e.g. verifiers) detect issues but rarely provide executable reproductions.
- Many bugs are cross-contract and multi-step (e.g. reentrancy)
 - hand-crafting attacker contracts + call sequences is slow and error-prone.
- In Solidity smart contracts, reproduction is more challenging
 - requires an attacker contract and a clear exploit path/trace.


Why Reproducing Bugs?

- **Understanding & Fixing:** Generate attacker contracts to help developers understand vulnerabilities.
 - (Ballesteros et al., 2022), (Wu et al., 2024).
- **Detection & Forecasting:** Pair adversarial with vulnerable contracts to detect/anticipate attacks.
 - (Wang et al., 2024)
- **Vulnerability Discovery:** Generate exploit PoCs to validate the (LLM-based) bug findings.
 - (Gervais & Zhou, 2025), (Xiao et al., 2025)

Our primary goal is *Understanding*, but the generated contracts are also suitable for *Forecasting*

Our Solution: VeriExploit

- LLM-guided generation of attacker contracts.
- Bounded cross-contract verification (BCCV).
- Automatically validates exploits and reports executable steps.
- Generalizes across multiple vulnerability types.



Algorithm Overview

Given:

- Vulnerable and attacker contracts.
- A verifier that supports unbounded reasoning.

VeriExploit checks:

- If an attacker path reproduces the reported bug, we provably find it.
- Otherwise, we formally diagnose why: **unreachable** or **incorrect parameters**.
(args, caller, value, pre-state)

Key Numbers

Dataset: 100 samples from open-sourced contacts, covering **six** bug types.

85.6%

Overall success
(ESBMC)

Validated exploits across
all benchmarks

64.6%

Overall success
(SolCMC)

Performance on the same
dataset

11.8%

LLM-only baseline
Without verifier feedback

+20.6 pp

Self-reflection gain

From 65.0% to 85.6% (ESBMC)

90.9%

AdvScanner comparison

vs 88.5% (AdvScanner)



The University of Manchester

Thank you!