

VERIFICATION OPTIMIZATION THROUGH GLOBAL COMMON SUBEXPRESSION ELIMINATION

Menezes, R. S., Tihanyi, N. Jain, R., Levin, A., de Freitas,
R. Cordeiro, L.

OUR TOOL

OUR TOOL

- ESBMC is a software verifier capable of ensuring safety properties of programs.

OUR TOOL

- ESBMC is a software verifier capable of ensuring safety properties of programs.
- It is used by industrial partners for the verification of firmware.

OUR TOOL

- ESBMC is a software verifier capable of ensuring safety properties of programs.
- It is used by industrial partners for the verification of firmware.
- The approach consists in translating a bounded trace of the program into an SMT formula.

OUR TOOL

- ESBMC is a software verifier capable of ensuring safety properties of programs.
- It is used by industrial partners for the verification of firmware.
- The approach consists in translating a bounded trace of the program into an SMT formula.
- <https://github.com/esbmc/esbmc>

HOW IT WORKS

```
1 int main() {  
2     char a,b;  
3     if(a > 10)  
4         b *= 2;  
5     assert(a + b != 42);  
6 }
```

HOW IT WORKS

```
1 int main() {  
2     char a,b;  
3     if(a > 10)  
4         b *= 2;  
5     assert(a + b != 42);  
6 }
```


HOW IT WORKS

```
1 int main() {  
2     char a,b;  
3     if(a > 10)  
4         b *= 2;  
5     assert(a + b != 42);  
6 }
```

HOW IT WORKS

```
1 int main() {  
2     char a,b;  
3     if(a > 10)  
4         b *= 2;  
5     assert(a + b != 42);  
6 }
```

HOW IT WORKS

```
1 int main() {  
2     char a,b;  
3     if(a > 10)  
4         b *= 2;  
5     assert(a + b != 42);  
6 }
```

$$(a_0 \geq -128 \wedge a_0 \leq 127) \wedge (b_0 \geq -128 \wedge b_0 \leq 127)$$

HOW IT WORKS

```
1 int main() {  
2     char a,b;  
3     if(a > 10)  
4         b *= 2;  
5     assert(a + b != 42);  
6 }
```

$$(a_0 \geq -128 \wedge a_0 \leq 127) \wedge (b_0 \geq -128 \wedge b_0 \leq 127) \\ \wedge (b_1 = b_0 \times 2)$$

HOW IT WORKS

```
1 int main() {  
2     char a,b;  
3     if(a > 10)  
4         b *= 2;  
5     assert(a + b != 42);  
6 }
```

$$(a_0 \geq -128 \wedge a_0 \leq 127) \wedge (b_0 \geq -128 \wedge b_0 \leq 127) \\ \wedge (b_1 = b_0 \times 2) \wedge [b_2 = \phi(a_0 > 10 \Rightarrow b_1, a_0 \leq 10 \Rightarrow b_0)]$$

HOW IT WORKS

```
1 int main() {  
2     char a,b;  
3     if(a > 10)  
4         b *= 2;  
5     assert(a + b != 42);  
6 }
```

$$\begin{aligned} & (a_0 \geq -128 \wedge a_0 \leq 127) \wedge (b_0 \geq -128 \wedge b_0 \leq 127) \\ & \wedge (b_1 = b_0 \times 2) \wedge [b_2 = \phi(a_0 > 10 \Rightarrow b_1, a_0 \leq 10 \Rightarrow b_0)] \\ & \wedge (a_0 + b_2 = 42) \end{aligned}$$

WHAT MAKES IT DIFFICULT

WHAT MAKES IT DIFFICULT

- Tracking pointers.

WHAT MAKES IT DIFFICULT

- Tracking pointers.
- We compute points-to by evaluating the symbolic expressions and recursively matching addresses.

WHAT MAKES IT DIFFICULT

- Tracking pointers.
- We compute points-to by evaluating the symbolic expressions and recursively matching addresses.
- This can be even trickier pointers with metadata.

WHAT MAKES IT DIFFICULT

- Tracking pointers.
- We compute points-to by evaluating the symbolic expressions and recursively matching addresses.
- This can be even trickier pointers with metadata.

```
1 struct my_ptr {  
2     int64_t size : 8;  
3     int64_t addr : 56;  
4 };
```

WHAT MAKES IT DIFFICULT

```
1 int foo() {
2     struct my_ptr ptr;
3     if(*) {
4         ptr.addr = malloc(120);
5         ptr.size = 120;
6     }
7     else {
8         ptr.addr = malloc(150);
9         ptr.size = 150;
10    }
11    int64_t my_addr = *((int64_t*)&ptr);
12    void *the_addr = (void*) (my_addr & 0x0FFFFFFF);
13    ...
14 }
```

WHAT MAKES IT DIFFICULT

```
1 int foo() {
2     struct my_ptr ptr;
3     if(*) {
4         ptr.addr = malloc(120);
5         ptr.size = 120;
6     }
7     else {
8         ptr.addr = malloc(150);
9         ptr.size = 150;
10    }
11    int64_t my_addr = *((int64_t*)&ptr);
12    void *the_addr = (void*) (my_addr & 0x0FFFFFFF);
13    ...
14 }
```

WHAT MAKES IT DIFFICULT

```
1 int foo() {
2     struct my_ptr ptr;
3     if(*) {
4         ptr.addr = malloc(120);
5         ptr.size = 120;
6     }
7     else {
8         ptr.addr = malloc(150);
9         ptr.size = 150;
10    }
11    int64_t my_addr = *((int64_t*)&ptr);
12    void *the_addr = (void*) (my_addr & 0x0FFFFFFF);
13    ...
14 }
```

WHAT MAKES IT DIFFICULT

```
1 int foo() {
2     struct my_ptr ptr;
3     if(*) {
4         ptr.addr = malloc(120);
5         ptr.size = 120;
6     }
7     else {
8         ptr.addr = malloc(150);
9         ptr.size = 150;
10    }
11    int64_t my_addr = *((int64_t*)&ptr);
12    void *the_addr = (void*) (my_addr & 0x0FFFFFFF);
13    ...
14 }
```

OPTIMIZING

OPTIMIZING

- ESBMC relies on multiple checks and analysis to compute pointer aliases.

OPTIMIZING

- ESBMC relies on multiple checks and analysis to compute pointer aliases.
- We can accelerate it and other operations by applying compiler optimizations.

EXAMPLE

```
1 void write(table *tbl, unsigned EntryIndex) {  
2     tbl->Map[EntryIndex].Auxiliary.Flags &= 1;  
3     tbl->Map[EntryIndex].Wc++;  
4     tbl->Map[EntryIndex].V = 42;  
5 }
```

EXAMPLE

```
1 void write(table *tbl, unsigned EntryIndex) {  
2     tbl->Map[EntryIndex].Auxiliary.Flags &= 1;  
3     tbl->Map[EntryIndex].Wc++;  
4     tbl->Map[EntryIndex].V = 42;  
5 }
```

EXAMPLE

```
1 void write(table *tbl, unsigned EntryIndex) {  
2     tbl->Map[EntryIndex].Auxiliary.Flags &= 1;  
3     tbl->Map[EntryIndex].Wc++;  
4     tbl->Map[EntryIndex].V = 42;  
5 }
```

- The repetitive dereferences makes ESBMC compute the same symbolic pointer operations over and over.

EXAMPLE

```
1 void write(table *tbl, unsigned EntryIndex) {  
2     tbl->Map[EntryIndex].Auxiliary.Flags &= 1;  
3     tbl->Map[EntryIndex].Wc++;  
4     tbl->Map[EntryIndex].V = 42;  
5 }
```

- The repetitive dereferences makes ESBMC compute the same symbolic pointer operations over and over.
- Invoking the `write` function over 1000 entries results in a 160s symbolic execution and memory exhaustion at solving (after 10min).

EXAMPLE

```
1 void write(table *tbl, unsigned EntryIndex) {  
2     entry *pEntry = &(tbl->Map[EntryIndex]);  
3     pEntry->Auxiliary.Flags &= 1;  
4     pEntry->Wc++;  
5     pEntry->V = 42;  
6 }
```

EXAMPLE

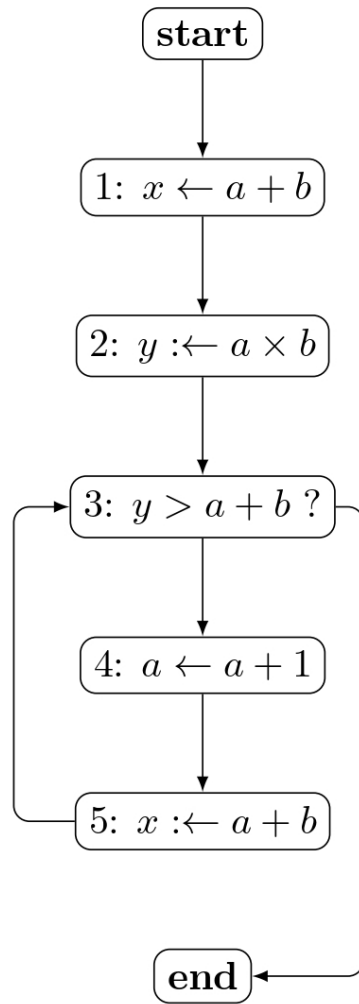
```
1 void write(table *tbl, unsigned EntryIndex) {  
2     entry *pEntry = &(tbl->Map[EntryIndex]);  
3     pEntry->Auxiliary.Flags &= 1;  
4     pEntry->Wc++;  
5     pEntry->V = 42;  
6 }
```


EXAMPLE

```
1 void write(table *tbl, unsigned EntryIndex) {  
2     entry *pEntry = &(tbl->Map[EntryIndex]);  
3     pEntry->Auxiliary.Flags &= 1;  
4     pEntry->Wc++;  
5     pEntry->V = 42;  
6 }
```

- By constructing an alias the analysis now takes less than 1s.

AVAILABLE EXPRESSIONS



Domain: $\{a + b, a \times b, a + 1\}$

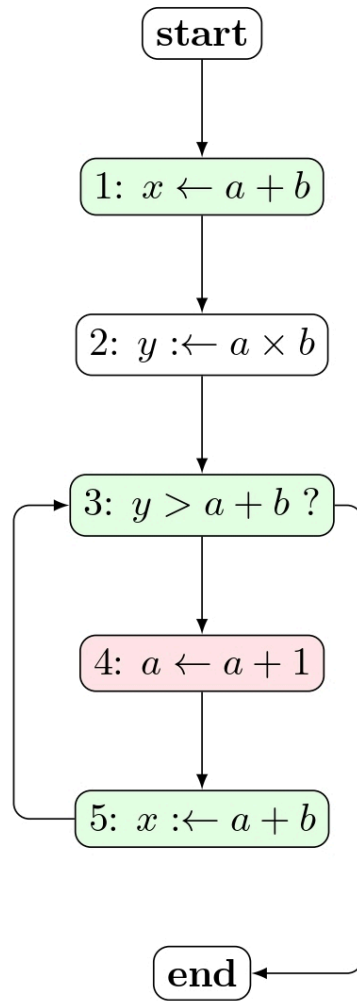
Vertex	Gen(v)	Kill(v)
1	$\{a + b\}$	\emptyset
2	$\{a \times b\}$	\emptyset
3	$\{a + b\}$	\emptyset
4	\emptyset	$\{a + b, a \times b, a + 1\}$
5	$\{a + b\}$	\emptyset

$$\text{In}(v) = \begin{cases} \mathbf{I}, v = \text{start} \\ \bigcap_{x \in \text{pred}(v)} \text{Out}(x) \end{cases}$$

$$\text{Out}(v) = \text{In}(v) \setminus \text{Kill}(v) \cup \text{Gen}(v)$$

Vertex	In(v)	Out(v)
1	\emptyset	$\{a + b\}$
2	$\{a + b\}$	$\{a + b, a \times b\}$
3	$\{a + b\}$	$\{a + b\}$
4	$\{a + b\}$	\emptyset
5	\emptyset	$\{a + b\}$

AVAILABLE EXPRESSIONS



Domain: $\{a + b, a \times b, a + 1\}$

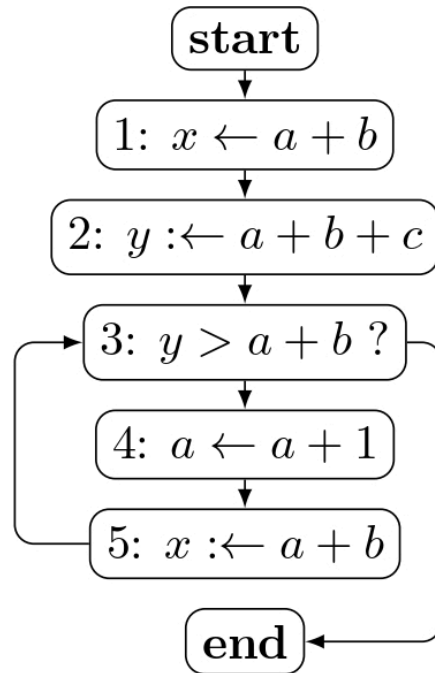
Vertex	Gen(v)	Kill(v)
1	$\{a + b\}$	\emptyset
2	$\{a \times b\}$	\emptyset
3	$\{a + b\}$	\emptyset
4	\emptyset	$\{a + b, a \times b, a + 1\}$
5	$\{a + b\}$	\emptyset

$$\text{In}(v) = \begin{cases} \mathbf{I}, v = \text{start} \\ \bigcap_{x \in \text{pred}(v)} \text{Out}(x) \end{cases}$$

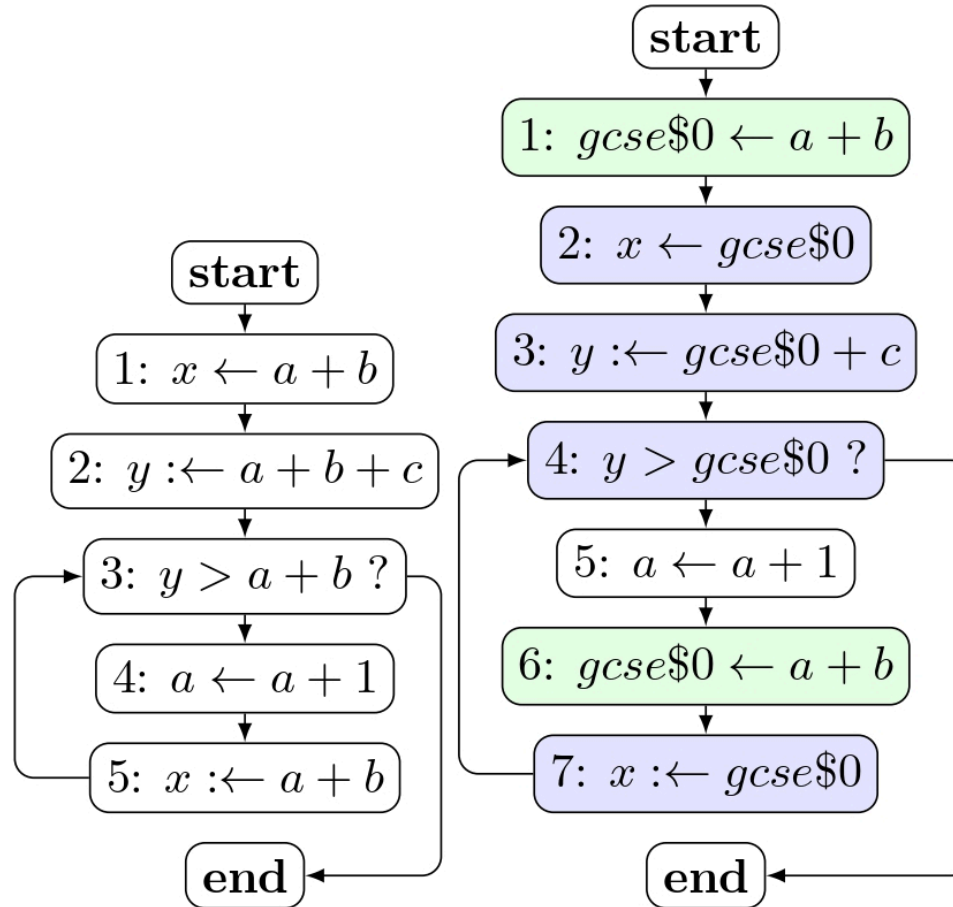
$$\text{Out}(v) = \text{In}(v) \setminus \text{Kill}(v) \cup \text{Gen}(v)$$

Vertex	In(v)	Out(v)
1	\emptyset	$\{a + b\}$
2	$\{a + b\}$	$\{a + b, a \times b\}$
3	$\{a + b\}$	$\{a + b\}$
4	$\{a + b\}$	\emptyset
5	\emptyset	$\{a + b\}$

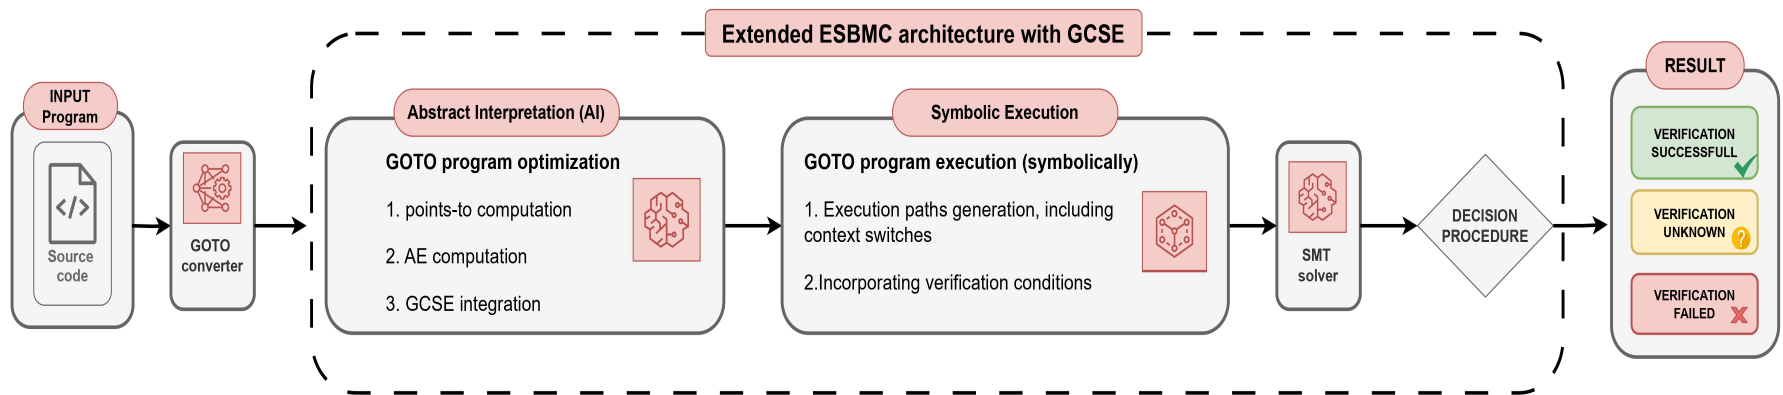
GCSE



GCSE



VO-GCSE



RESULTS

RESULTS

- We tested our approach on SV-Comp benchmarks for Reachsafety and Memory categories.

RESULTS

- We tested our approach on SV-Comp benchmarks for Reachsafety and Memory categories.
- VO-GCSE delivers significant performance improvements in verification time— up to 52%—in memory-related verification tasks.

RESULTS

- We tested our approach on SV-Comp benchmarks for Reachsafety and Memory categories.
- VO-GCSE delivers significant performance improvements in verification time— up to 52%—in memory-related verification tasks.
- VO-GCSE struggles in categories such as Hardware and ECA. Mainly due to the overhead of the Abstract Interpreter.

NEXT STEPS

NEXT STEPS

- We aim to extend ESBMC to be able to use modern compiler techniques, including: phi-nodes, IFDS, data-flow.

NEXT STEPS

- We aim to extend ESBMC to be able to use modern compiler techniques, including: phi-nodes, IFDS, data-flow.
- Improve the Abstract Interpreter.

NEXT STEPS

- We aim to extend ESBMC to be able to use modern compiler techniques, including: phi-nodes, IFDS, data-flow.
- Improve the Abstract Interpreter.
- Improve the points-to analysis.

NEXT STEPS

- We aim to extend ESBMC to be able to use modern compiler techniques, including: phi-nodes, IFDS, data-flow.
- Improve the Abstract Interpreter.
- Improve the points-to analysis.
- Explore other compiler techniques that can be applied to software verification (e.g., loop transformations, strength weakening).

THANK YOU