

Collaborators/funders:

APT / FM / S3 Research Groups

ARM Centre of Excellence

Centre for Digital Trust and Society

PPGEE, PPGI – UFAM

Innovate UK, UKRI, EPSRC, and EU Horizon

Industrial partners (ARM, Ethereum, Intel, Motorola, TII, Zscaler)



The University of Manchester

Towards Building Trustworthy Software and AI Systems



Lucas C. Cordeiro
Department of Computer Science
lucas.cordeiro@manchester.ac.uk
<https://ssvlab.github.io/lucasccordeiro/>



Career Summary

1



BSc/MSc in
Engineering and
Lecturer

7

2



MSc in Embedded
Systems

3

SIEMENS
mobile

Configuration and
Build Manager

4

BenQ
SIEMENS

Feature Leader

5

NMP

Set-top Box
Software Engineer

6

Microsoft



PhD in Computer
Science

8



Postdoctoral
Researcher

9

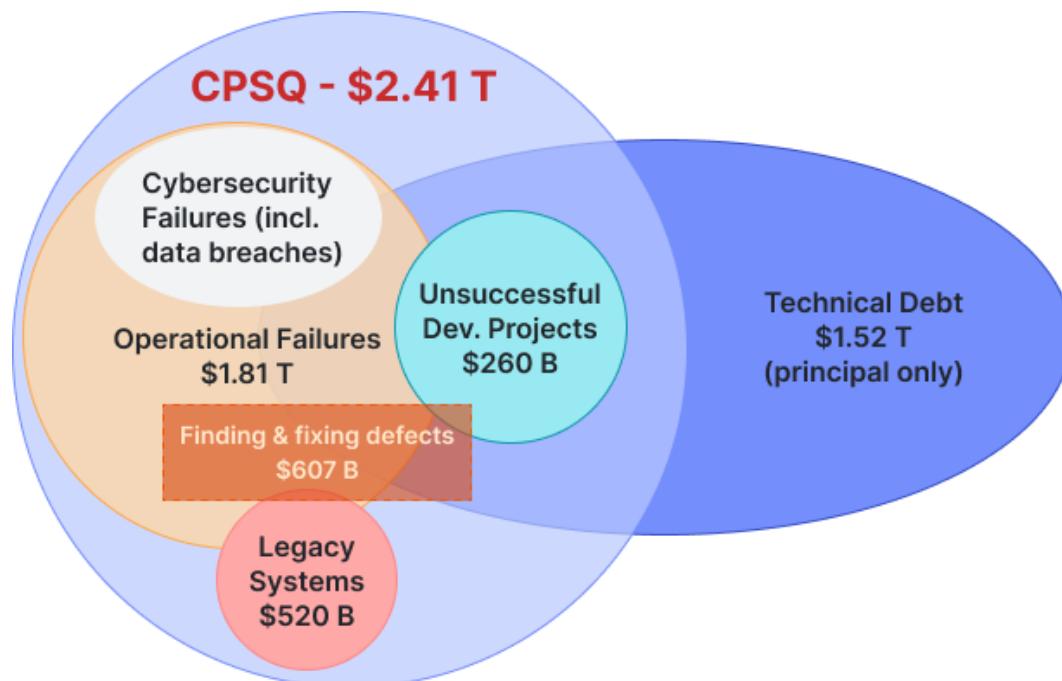
MANCHESTER
1824

The University of Manchester

Professor of
Computer Science

How much could software errors cost your business?

Poor software quality cost US companies \$2.41 trillion in 2022, while the accumulated software Technical Debt (TD) has grown to ~\$1.52 trillion



TD relies on temporary easy-to-implement solutions to achieve short-term results at the expense of efficiency in the long run

The cost of poor software quality in the US: A 2022 Report

How secure is AI-generated Code: A Large-Scale Comparison of Large Language Models

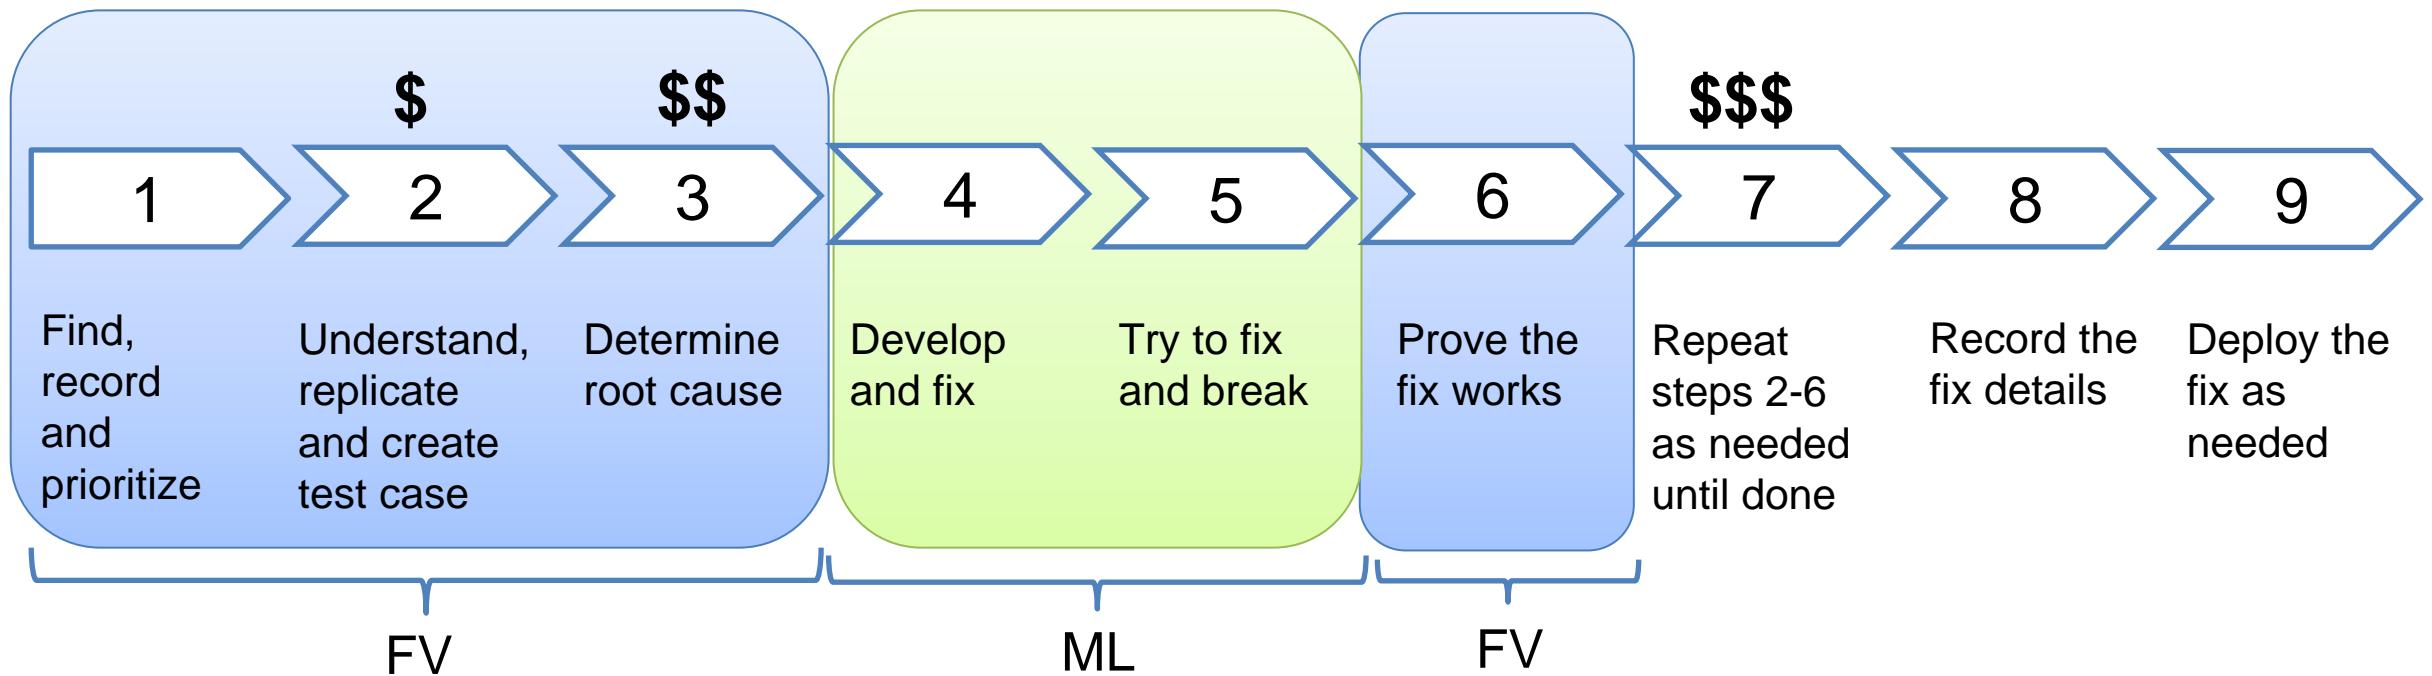
Category	Avg Prop. Viol. per Line	Rank	\mathcal{VS}	Rank	\mathcal{VF}	\mathcal{VU} (Timeout)	Avg Prop. Viol. per File
GPT-4o-mini	0.0165	3	4.23%	2	57.14%	36.77%	3.40
Llama2-13B	0.0234	2	12.36%	1	51.30%	31.78%	3.62
Mistral-7B	0.0254	7	8.36%	4	62.08%	25.88%	3.07
CodeLlama-13B	0.0260	1	15.48%	3	52.71%	29.52%	4.13
Falcon-180B	0.0291	8	6.48%	5	62.07%	28.67%	3.38
GPT-3.5-turbo	0.0295	6	7.29%	7	65.07%	26.09%	4.42
Gemini Pro 1.0	0.0305	5	9.49%	6	63.91%	24.13%	4.70
Gemma-7B	0.0437	4	11.62%	8	67.01%	16.30%	4.20

Legend:

\mathcal{VS} : 0.0234 Verification Success; \mathcal{VF} : Verification Failed; \mathcal{VU} : Verification Unknown (Timeout).

Best performance in a category is highlighted with bold and/or Rank.

Find, Understand and Fix Bugs



“A significant percentage (50%+) of a software project’s cost today is not spent on the creativity activity of software construction but rather on the corrective activity of debugging and fixing errors”

The cost of poor software quality
in the US: A 2022 Report

What do you think, Julia?



That's not good,
that's not nice,
that's not funny. It
is scary!

Objective of this talk

Discuss automated testing, verification, and synthesis to establish a foundation for building trustworthy software and AI systems

- Introduce a **logic-based automated reasoning platform** to find and fix **software defects**
- Explain **testing, verification, and synthesis** techniques to build **trustworthy software and AI systems**
- Develop an **automated reasoning system** for **safeguarding software and AI systems** against vulnerabilities in an increasingly digital and interconnected world

Research Questions

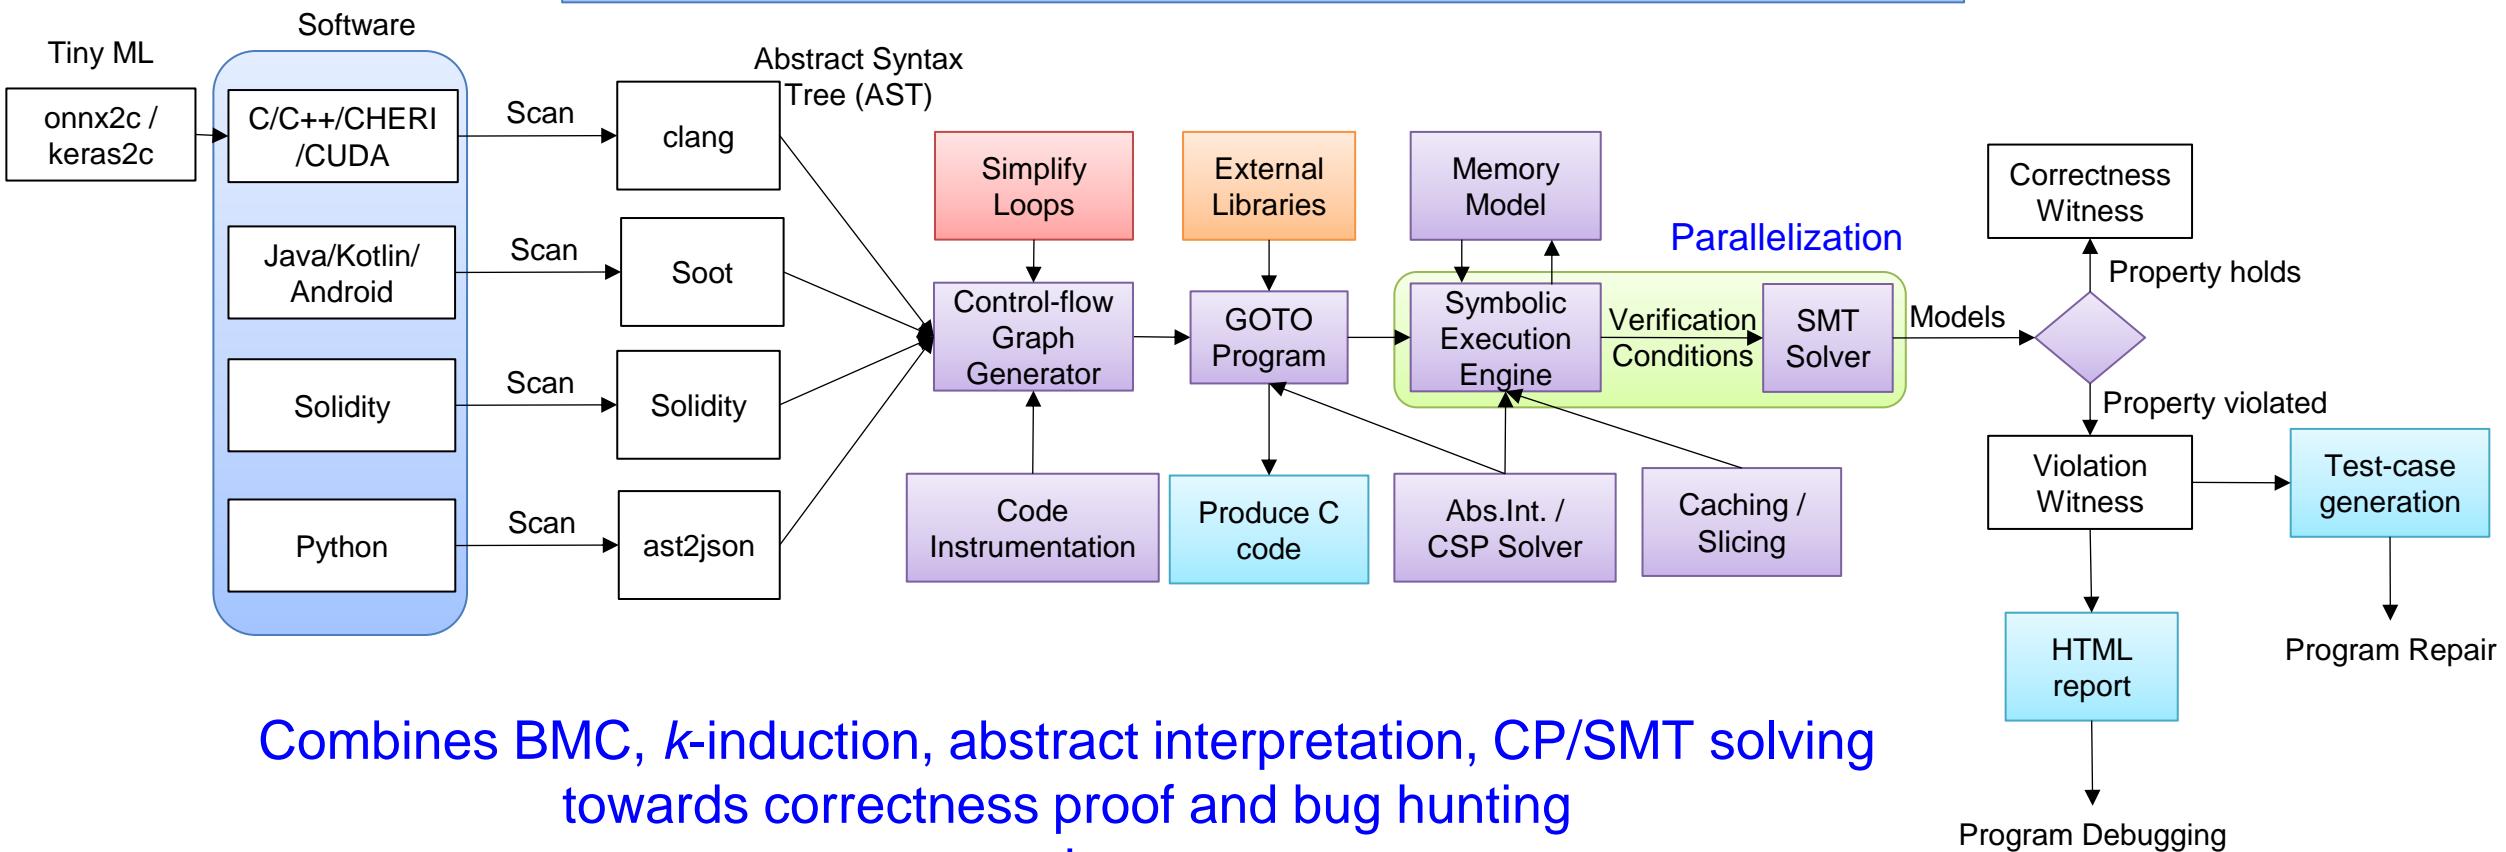
Given a **program** and a **specification**, can we automatically **verify** that the **program performs as specified**?

Can we leverage **program analysis/repair** to **discover and fix** more **software vulnerabilities** than existing state-of-the-art approaches?

Can we **improve engineers' productivity** to **find, understand, and fix software vulnerabilities**?

ESBMC: A Logic-based Verification Platform

Logic-based automated verification
for checking **safety** and **liveness**
properties in **AI** and **software systems**



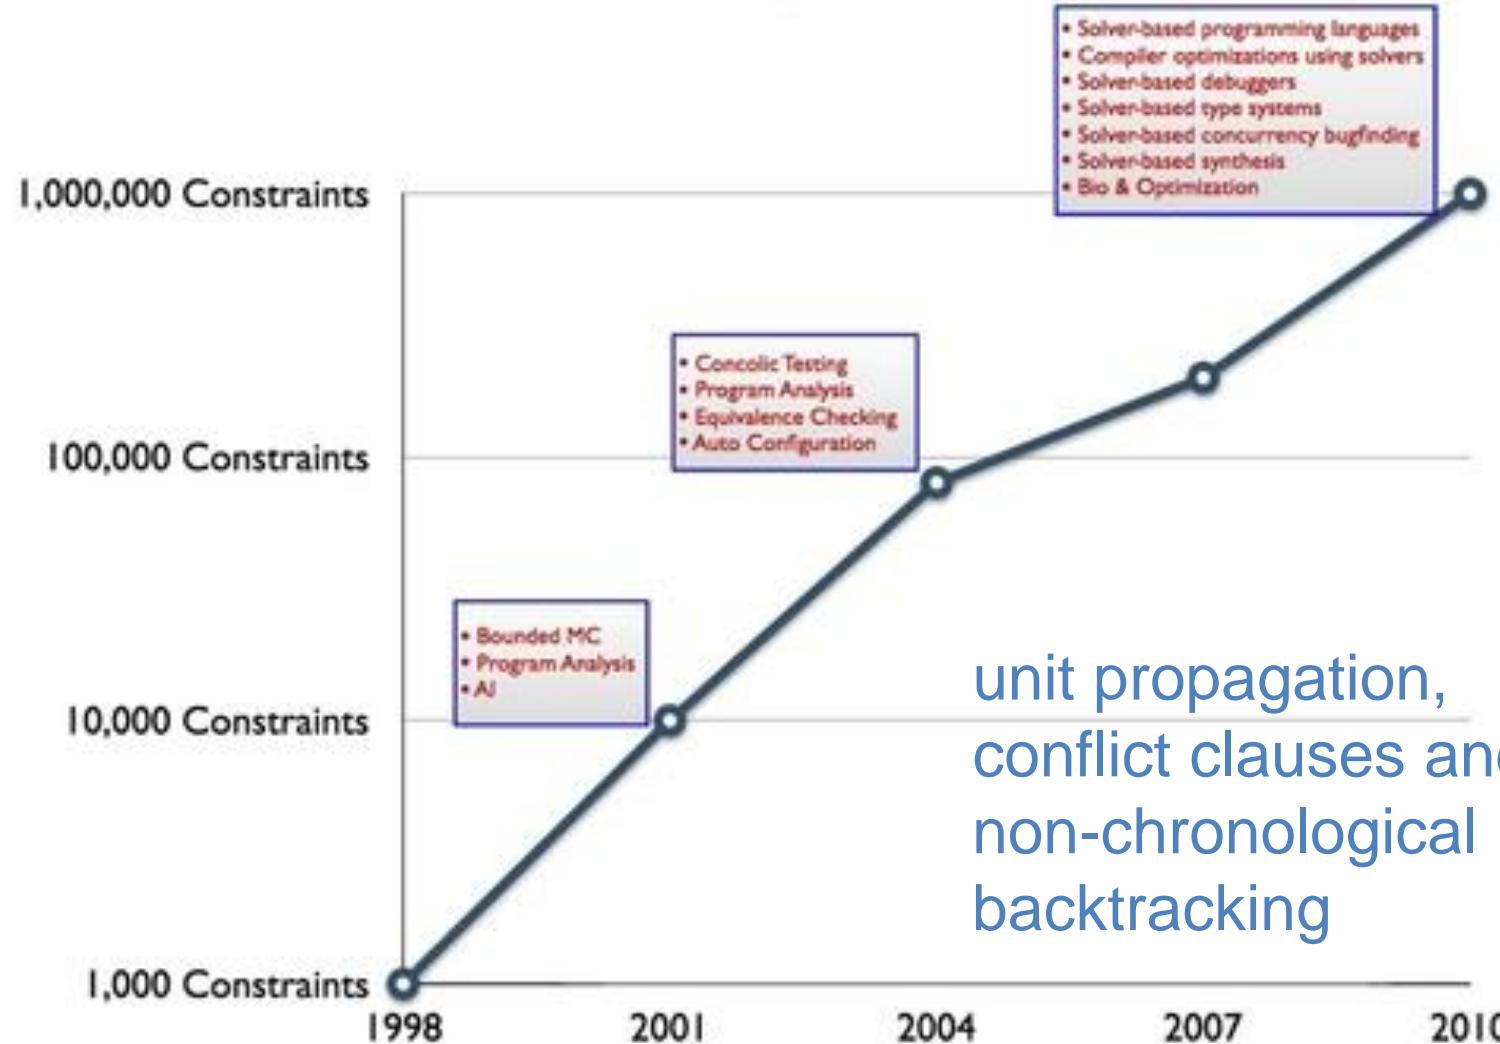
Agenda

- Automated program verification and testing to ensure conformance to specifications
- Advancing vulnerability detection to enhance software security beyond the state-of-the-art
- Automated program repair to improve the detection, comprehension, and fixing of software defects
- Summary and vision for future exploration

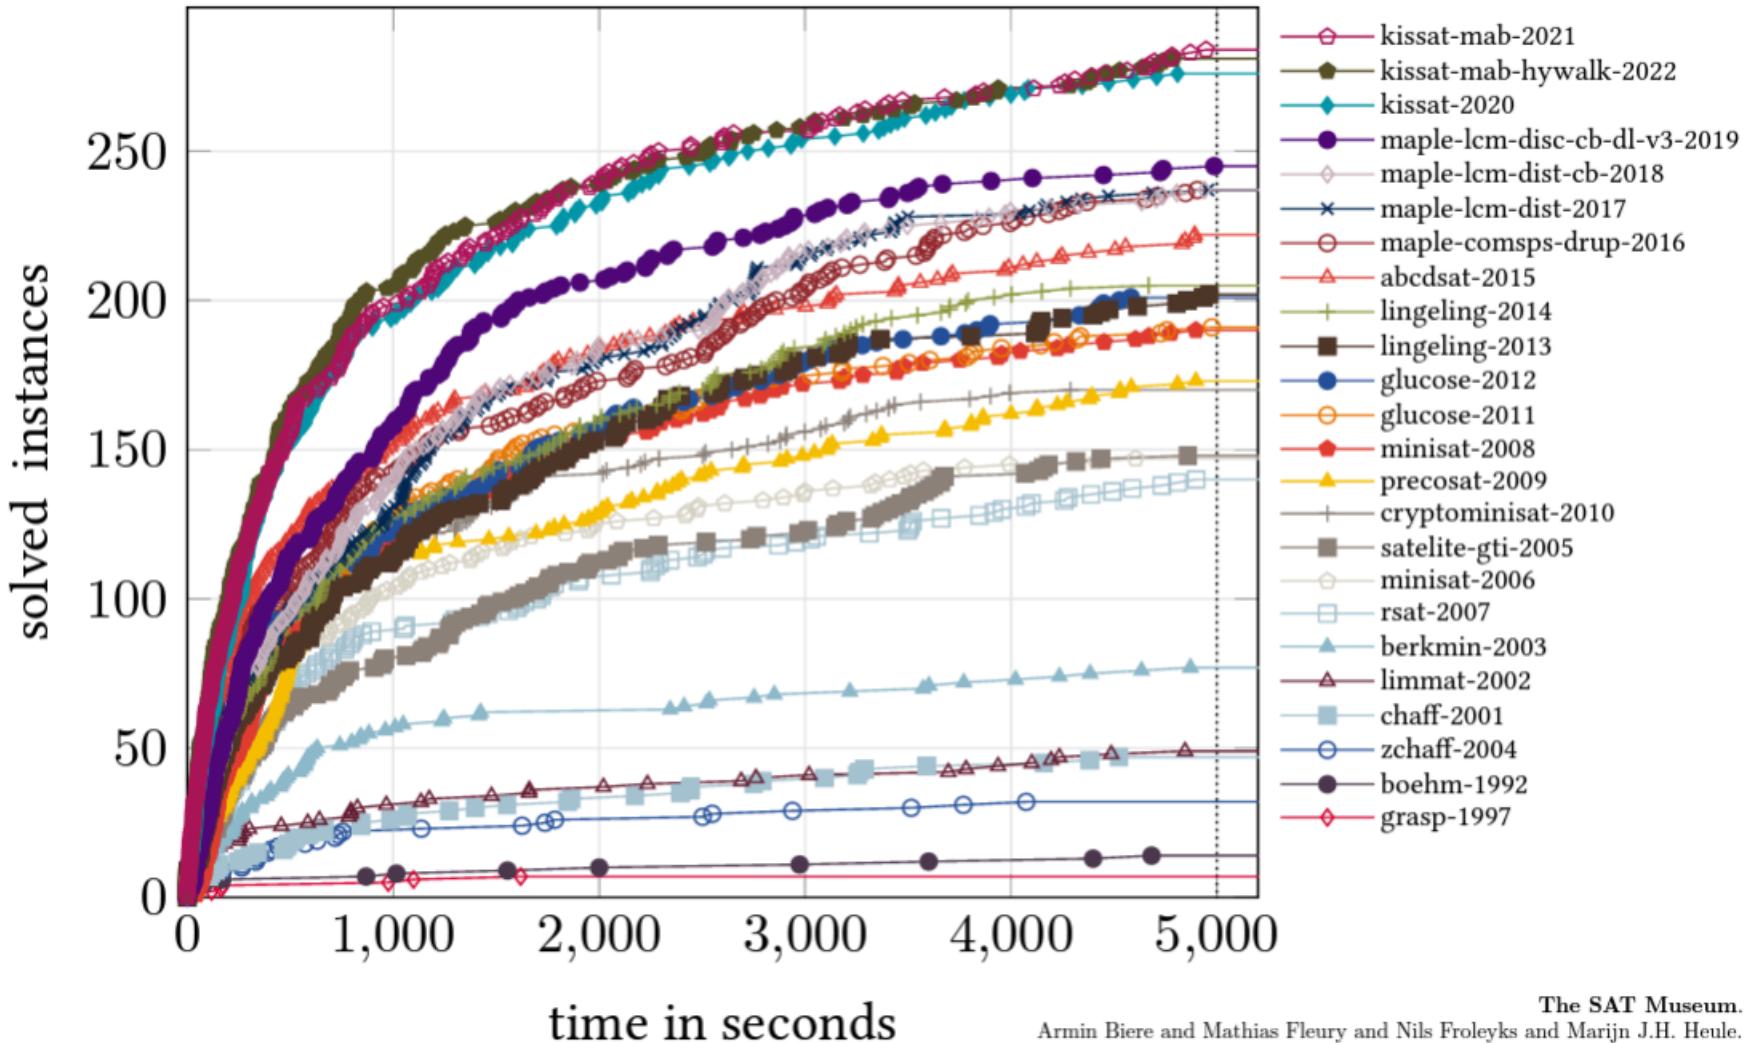
SAT solving as enabling technology



SAT/SMT Solver Research Story A 1000x Improvement



SAT Competition All Time Winners on SAT Competition 2022 Benchmarks



The SAT Museum.

Armin Biere and Mathias Fleury and Nils Froleyks and Marijn J.H. Heule.
In *Proceedings 14th International Workshop on Pragmatics of SAT (POS'23)*,

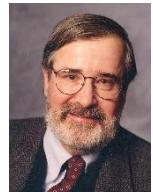
vol. 3545, CEUR Workshop Proceedings, pages 72-87, CEUR-WS.org 2023.

[paper - bibtex - data - zenodo - ceur - workshop - proceedings]

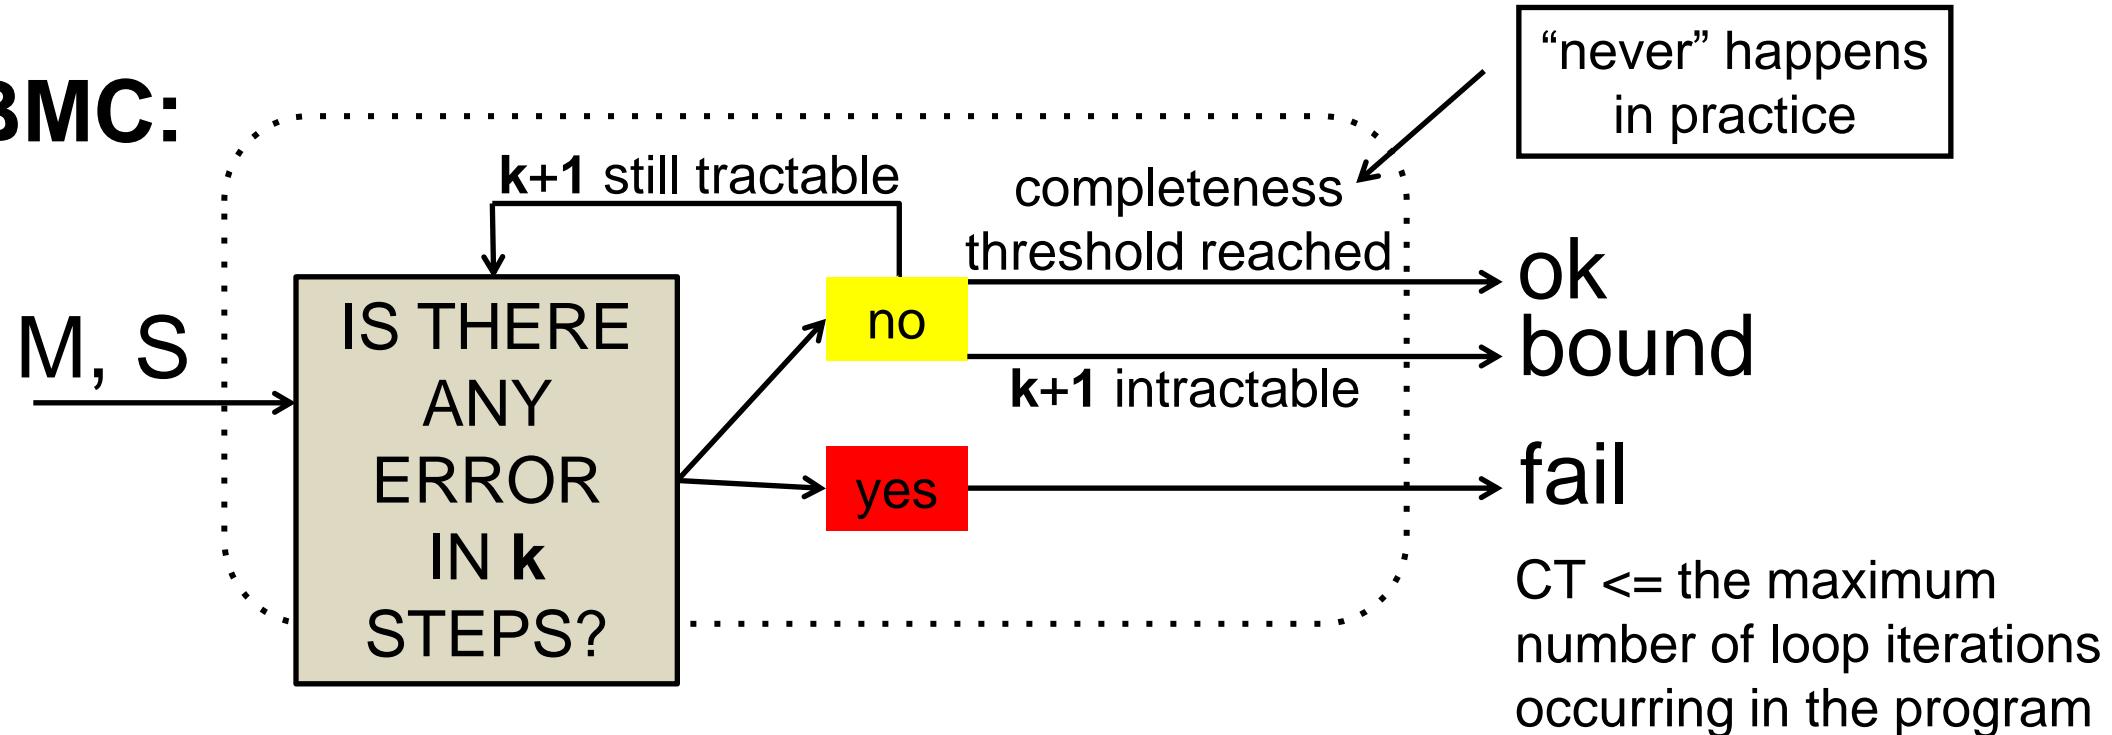
<https://cca.informatik.uni-freiburg.de/satmuseum>

<https://cca.informatik.uni-freiburg.de/satmuseum/>

Bounded Model Checking (BMC)



BMC:



Can the given property fail in k -steps?

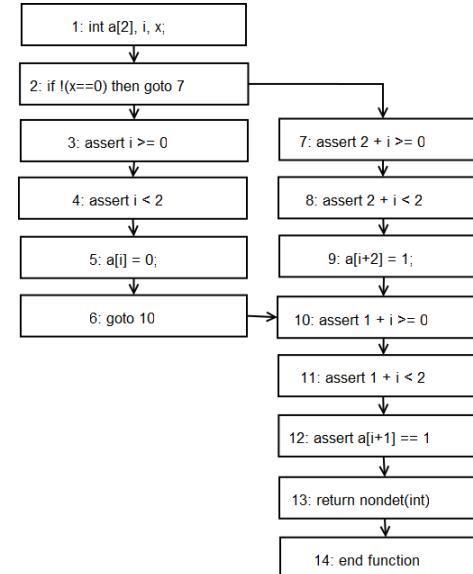
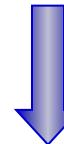
$$I(S_0) \wedge T(S_0, S_1) \wedge \dots \wedge T(S_{k-1}, S_k) \wedge (\neg P(S_0) \vee \dots \vee \neg P(S_k))$$

Property fails in some step

Software BMC

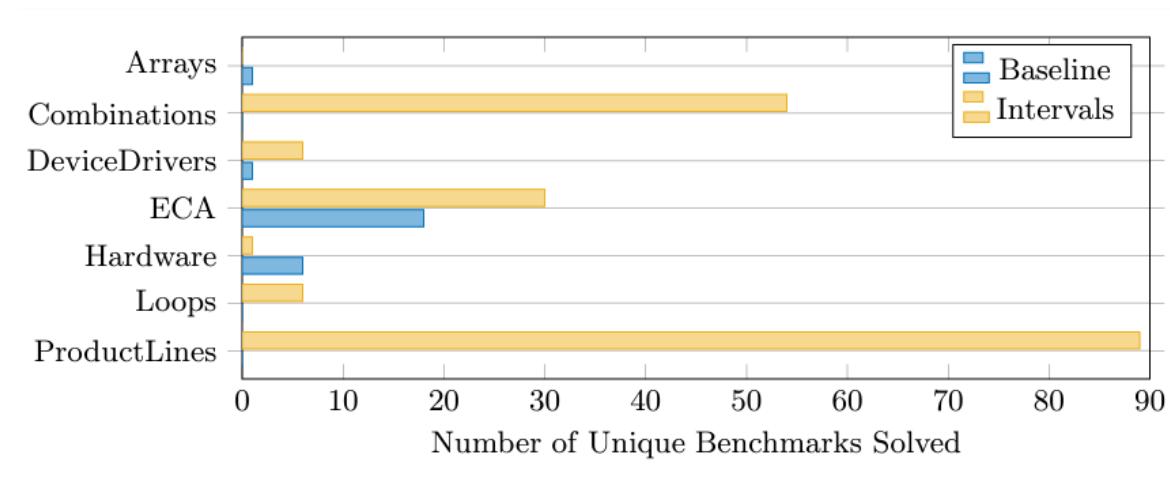
- program modeled as a state transition system
 - *state*: pc and program variables
 - derived from control-flow graph

```
int main() {  
    int a[2], i, x;  
    if (x==0)  
        a[i]=0;  
    else  
        a[i+2]=1;  
    assert(a[i+1]==1);  
}
```



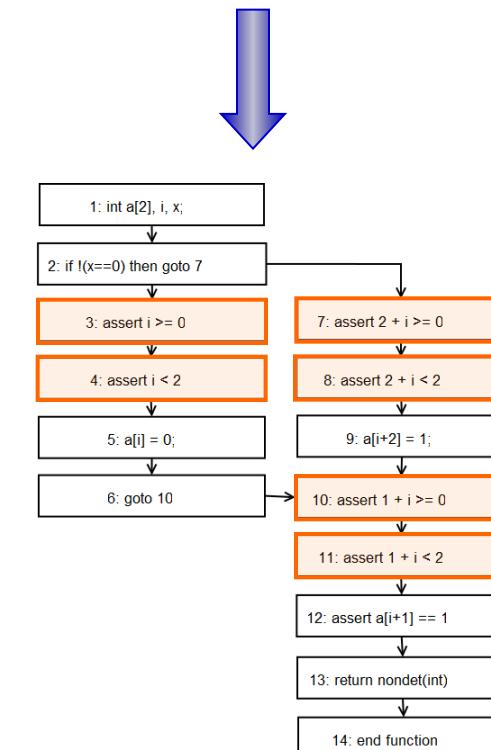
Software BMC

- program modeled as a state transition system
 - *state*: pc and program variables
 - derived from control-flow graph
 - added assumptions/safety properties as extra nodes



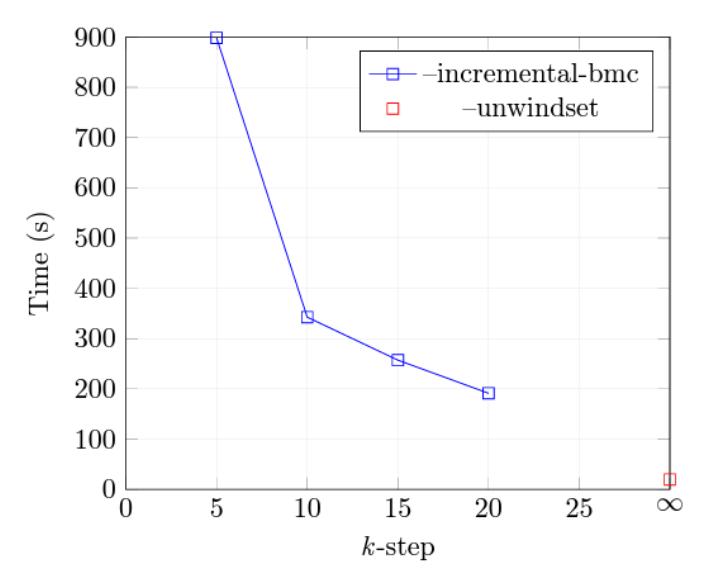
Menezes, R., Manino, E., Shmarov, F., Aldughaim, M., de Freitas, R., Lucas C. Cordeiro: Interval Analysis in Industrial-Scale BMC Software Verifiers: A Case Study.

```
int main() {  
    int a[2], i, x;  
    if (x==0)  
        a[i]=0;  
    else  
        a[i+2]=1;  
    assert(a[i+1]==1);  
}
```



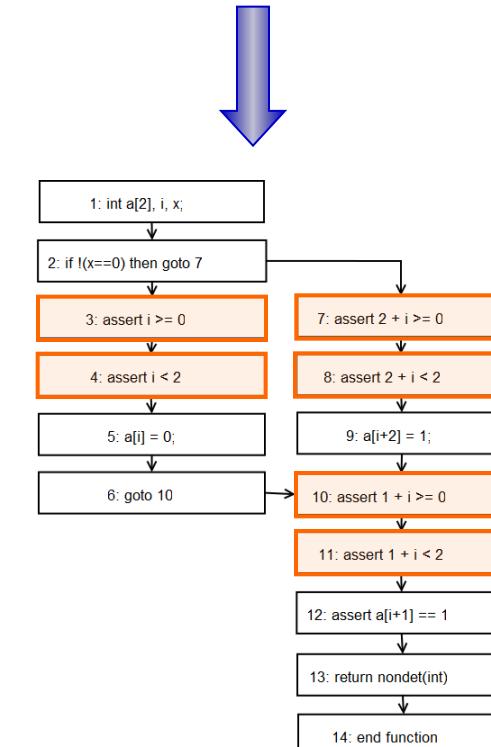
Software BMC

- program modeled as a state transition system
 - *state*: pc and program variables
 - derived from control-flow graph
 - added assumptions/safety properties as extra nodes
- program unfolded up to given bounds



Wu T., Xiong, S., Manino, E., Stockwell, G., Cordeiro, L.:
Verifying components of Arm(R) Confidential Computing
Architecture with ESBMC. SAS 2024 (to appear)

```
int main() {  
    int a[2], i, x;  
    if (x==0)  
        a[i]=0;  
    else  
        a[i+2]=1;  
    assert(a[i+1]==1);  
}
```

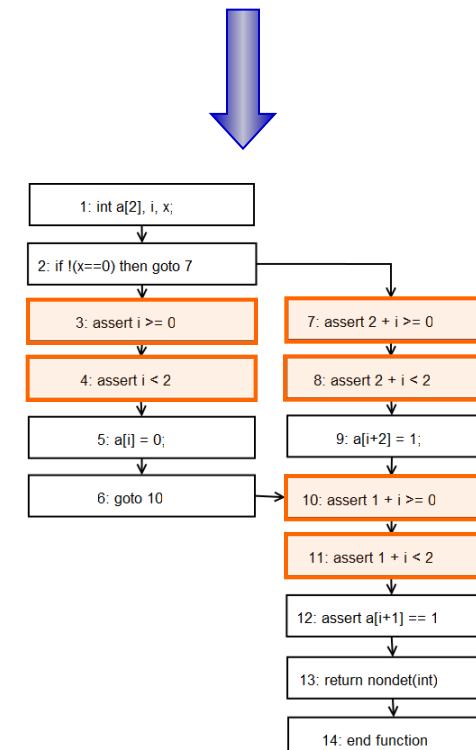


Software BMC

- program modeled as a state transition system
 - *state*: pc and program variables
 - derived from control-flow graph
 - added assumptions/safety properties as extra nodes
- program unfolded up to given bounds
- unfolded program optimized to reduce blow-up
 - constant propagation/slicing
 - forward substitutions/caching
 - unreachable code/pointer analysis

} crucial

```
int main() {  
    int a[2], i, x;  
    if (x==0)  
        a[i]=0;  
    else  
        a[i+2]=1;  
    assert(a[i+1]==1);  
}
```

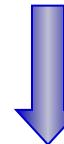


Software BMC

- program modeled as a state transition system
 - *state*: pc and program variables
 - derived from control-flow graph
 - added assumptions/safety properties as extra nodes
- program unfolded up to given bounds
- unfolded program optimized to reduce blow-up
 - constant propagation/slicing
 - forward substitutions/caching
 - unreachable code/pointer analysis
- front-end converts unrolled and **optimized program into SSA**

} crucial

```
int main() {  
    int a[2], i, x;  
    if (x==0)  
        a[i]=0;  
    else  
        a[i+2]=1;  
    assert(a[i+1]==1);  
}
```

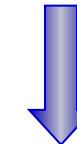


$g_1 = x_1 == 0$
 $a_1 = a_0 \text{ WITH } [i_0 := 0]$
 $a_2 = a_0$
 $a_3 = a_2 \text{ WITH } [2+i_0 := 1]$
 $a_4 = g_1 ? a_1 : a_3$
 $t_1 = a_4 [1+i_0] == 1$

Software BMC

- program modeled as a state transition system
 - *state*: pc and program variables
 - derived from control-flow graph
 - added assumptions/safety properties as extra nodes
- program unfolded up to given bounds
- unfolded program optimized to reduce blow-up
 - constant propagation/slicing
 - forward substitutions/caching
 - unreachable code/pointer analysis
- front-end converts unrolled and **optimized program into SSA**
- extraction of *constraints C* and *properties P*
 - specific to selected SMT solver, uses theories
- satisfiability check of $C \wedge \neg P$

```
int main() {  
    int a[2], i, x;  
    if (x==0)  
        a[i]=0;  
    else  
        a[i+2]=1;  
    assert(a[i+1]==1);  
}
```



$$C := \left[\begin{array}{l} g_1 := (x_1 = 0) \\ \wedge a_1 := \text{store}(a_0, i_0, 0) \\ \wedge a_2 := a_0 \\ \wedge a_3 := \text{store}(a_2, 2 + i_0, 1) \\ \wedge a_4 := \text{ite}(g_1, a_1, a_3) \end{array} \right]$$

$$P := \left[\begin{array}{l} i_0 \geq 0 \wedge i_0 < 2 \\ \wedge 2 + i_0 \geq 0 \wedge 2 + i_0 < 2 \\ \wedge 1 + i_0 \geq 0 \wedge 1 + i_0 < 2 \\ \wedge \text{select}(a_4, i_0 + 1) = 1 \end{array} \right]$$

Most Influential Paper Award at ASE 2023



Context-Bounded Model Checking in ESBMC

**Idea: iteratively generate all possible interleavings
and call the BMC procedure on each interleaving**

... combines

- **symbolic** model checking: on each individual interleaving
- **explicit state** model checking: explore all interleavings
 - bound the number of context switches allowed among threads

... implements

- **symbolic state hashing** (SHA1 hashes)
 - **monotonic partial order** reduction that combines dynamic POR with symbolic state space exploration

DISTINGUISHED PAPER AWARD

ICSE 2011

**The 33rd International Conference on
Software Engineering**

May 21-28, 2011

Waikiki, Honolulu, Hawaii

Presented to

Lucas Cordeiro and Bernd Fischer

For

**"Verifying Multi-threaded Software using
SMT-based Context-Bounded**

Competition on Software Verification (SV-COMP)

ControlFlowInteger

1. CPAchecker-ABE 1.0.10
2. CPAchecker-Memo 1.0.10
3. QARMC-HSF
4. ESBMC 1.17
5. LLBMC 0.9

DeviceDrivers

1. LLBMC 0.9
2. Predator
3. BLAST 2.7
4. SATabs 3.0
5. Wolverine 0.5c

DeviceDrivers64

1. BLAST 2.7
2. CPAchecker-Memo 1.0.10
3. SATabs 3.0
4. CPAchecker-ABE 1.0.10
5. Wolverine 0.5c

HeapManipulation

1. Predator
2. LLBMC 0.9
3. CPAchecker-ABE 1.0.10
3. CPAchecker-Memo 1.0.10
5. ESBMC 1.17

SystemC

1. ESBMC 1.17
2. SATabs 3.0
3. CPAchecker-ABE 1.0.10
4. CPAchecker-Memo 1.0.10
5. Wolverine 0.5c

Concurrency

1. ESBMC 1.17
2. SATabs 3.0
3. --
4. --
5. --

Overall

1. CPAchecker-Memo 1.0.10
2. CPAchecker-ABE 1.0.10
3. ESBMC 1.17
4. SATabs 3.0
5. BLAST 2.7

Induction-Based Verification for Software



k -induction checks loop-free programs...

- **base case (base_k)**: find a counter-example with up to k loop unwindings (plain BMC)
- **forward condition (fwd_k)**: check that P holds in all states reachable within k unwindings
- **inductive step (step_k)**: check that whenever P holds for k unwindings, it also holds after next unwinding
 - havoc variables
 - assume loop condition
 - run loop body (k times)
 - assume loop termination

⇒ iterative deepening if inconclusive

Gadelha, M., Ismail, H., Cordeiro, L.: Handling loops in bounded model checking of C programs via k -induction. Int. J. Softw. Tools Technol. Transf. 19(1): 97-114 (2017)

Automatic Invariant Generation



- Infer invariants based on **intervals** as abstract domain via a dependence graph
 - E.g., $a \leq x \leq b$ (integer and floating-point)
 - Inject intervals as assumptions and contract them via CSP
 - Remove unreachable states

Line	Interval for “a”	Restriction
4	$(-\infty, +\infty)$	None
6	$(-\infty, 100]$	$a \leq 100$
7	$(100, +\infty)$	$a > 100$

```
1 int main()
2 {
3     int a = *;
4     while(a <= 100)
5         a++;
6     assert(a>10);
7     return 0;
8 }
```

k-Induction proof rule “hijacks” loop conditions to nondeterministic values, thus computing intervals become essential

***k*-Induction can prove the correctness of more programs when the invariant generation is enabled**

Menezes, et al.: ESBMC v7.4: Harnessing the Power of Intervals - (Competition Contribution). TACAS (3) 2024: 376-380

BMC of Software Using Interval Methods via Contractors



- 1) Analyze intervals and properties
 - Static Analysis / Abstract Interpretation
- 2) Convert the problem into a CSP
 - Variables, Domains and Constraints
- 3) Apply contractor to CSP
 - Forward-Backward Contractor
- 4) Apply reduced intervals back to the program

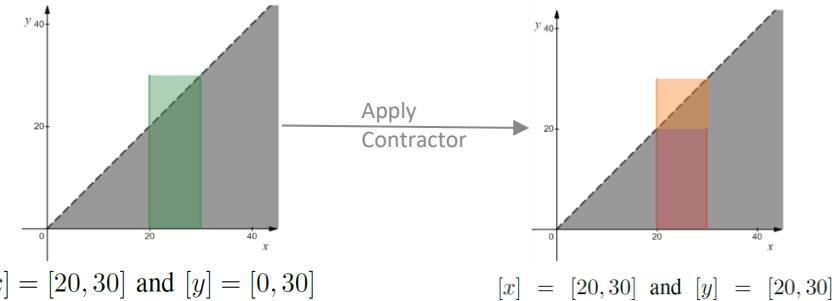
```
1 unsigned int x=nondet_uint();  
2 unsigned int y=nondet_uint();  
3 __ESBMC_assume(x >= 20 && x <= 30);  
4 __ESBMC_assume(y <= 30);  
5 assert(x >= y);  
  
    __ESBMC_assume(y <= 30 && y >= 20);
```

This **assumption** prunes our search space to the **orange area**

```
1 unsigned int x=nondet_uint();  
2 unsigned int y=nondet_uint();  
3 __ESBMC_assume(x >= 20 && x <= 30);  
4 __ESBMC_assume(y <= 30);  
5 assert(x >= y);
```

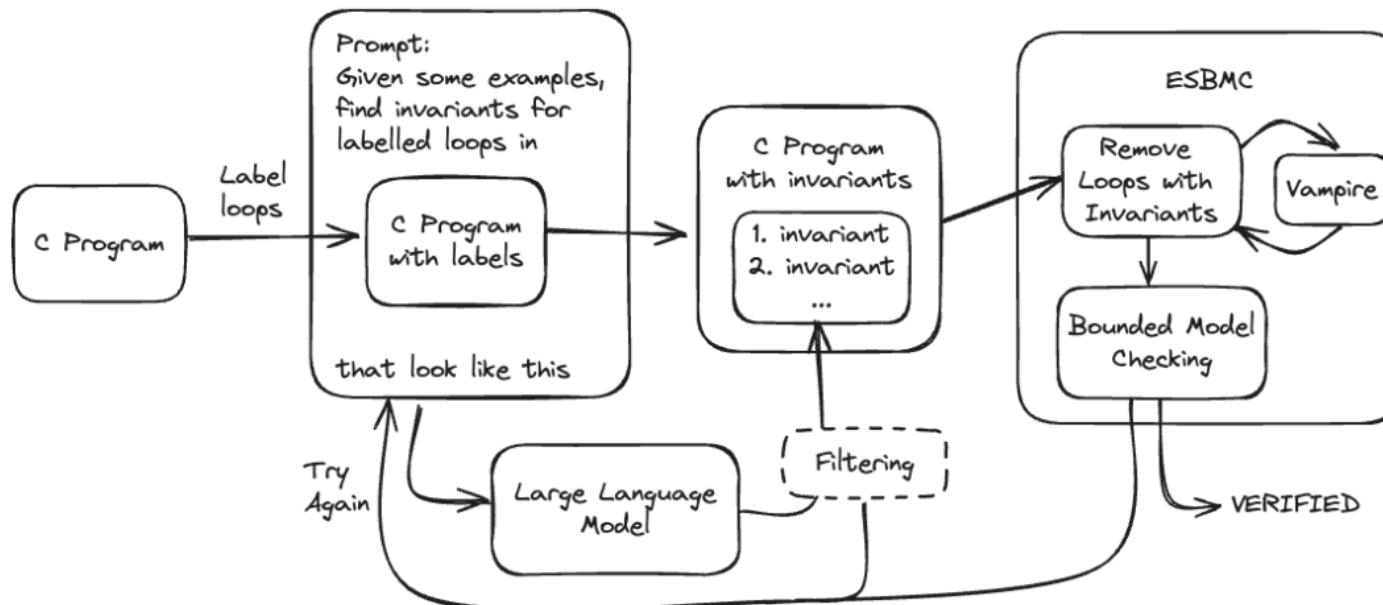
Domain: $[x] = [20, 30]$ and $[y] = [0, 30]$

Constraint: $y - x \leq 0$



$f(x) > 0$	$I = [0, \infty)$	
$f(x) = y - x$	$[f(x)_1] = I \cap [y_0] - [x_0]$	Forward-step
$x = y - f(x)$	$[x_1] = [x_0] \cap [y_0] - [f(x)_1]$	Backward-step
$y = f(x) + x$	$[y_1] = [y_0] \cap [f(x)_1] + [x_1]$	Backward-step

LLM-Generated Invariants for Bounded Model Checking Without Loop Unrolling



```
int main()
{
    int x = 0;
    int y = 50;
    __invariant(0 <= x && x <= 100);
    __invariant(x <= 50 && y == 50
               || x > 50 && y == x);
    while (x < 100) {
        x = x + 1;
        if(x > 50){
            y = y + 1;
        }
    }
    __VERIFIER_assert(y == 100 );
}
```

→ x = 0
y = 50
(0 <= x && x <= 100)
(x <= 50 && y == 50 || x > 50 && y == x)
F
assert y == 100
→ x = 0
y = 50
assume ((0 <= x && x <= 100) &&
(x <= 50 && y == 50 || x > 50 && y == x)
&& x >= 100)
↓
assert y == 100



Distinguished Paper Award

ASE 2024

IEEE/ACM International Conference on Automated Software Engineering

October 27 - November 1

Sacramento, California

Presented to

Muhammad A. A. Pirzada, Giles Reger, Ahmed Bhayat, Lucas C. Cordeiro

for

**LLM-Generated Invariants for Bounded Model
Checking Without Loop Unrolling**

Vladimir Filkov

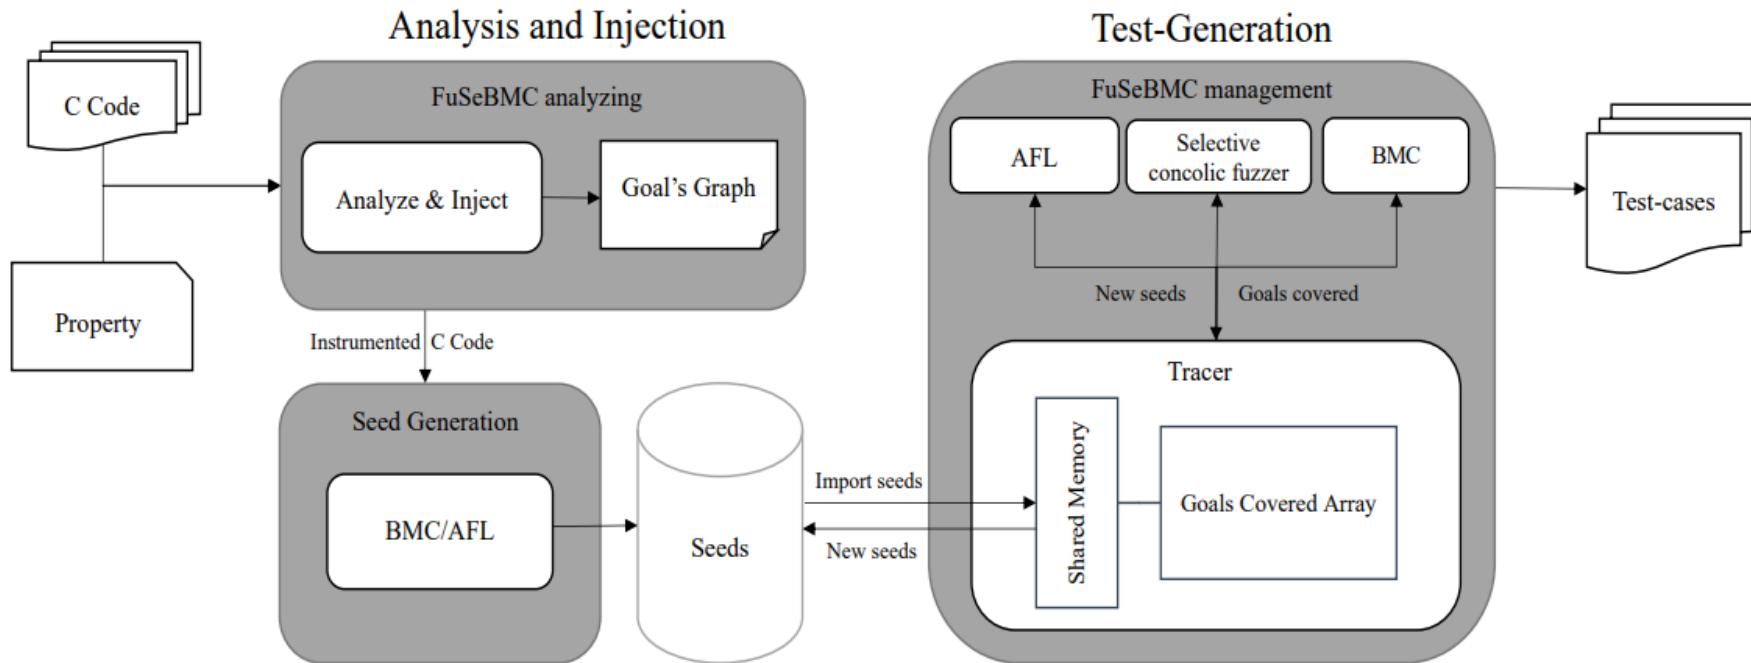
Vladimir Filkov
General Chair

Baishakhi Ray
Research PC Co-Chair

Minghui Zhou
Research PC Co-Chair

FuSeBMC v4 Framework

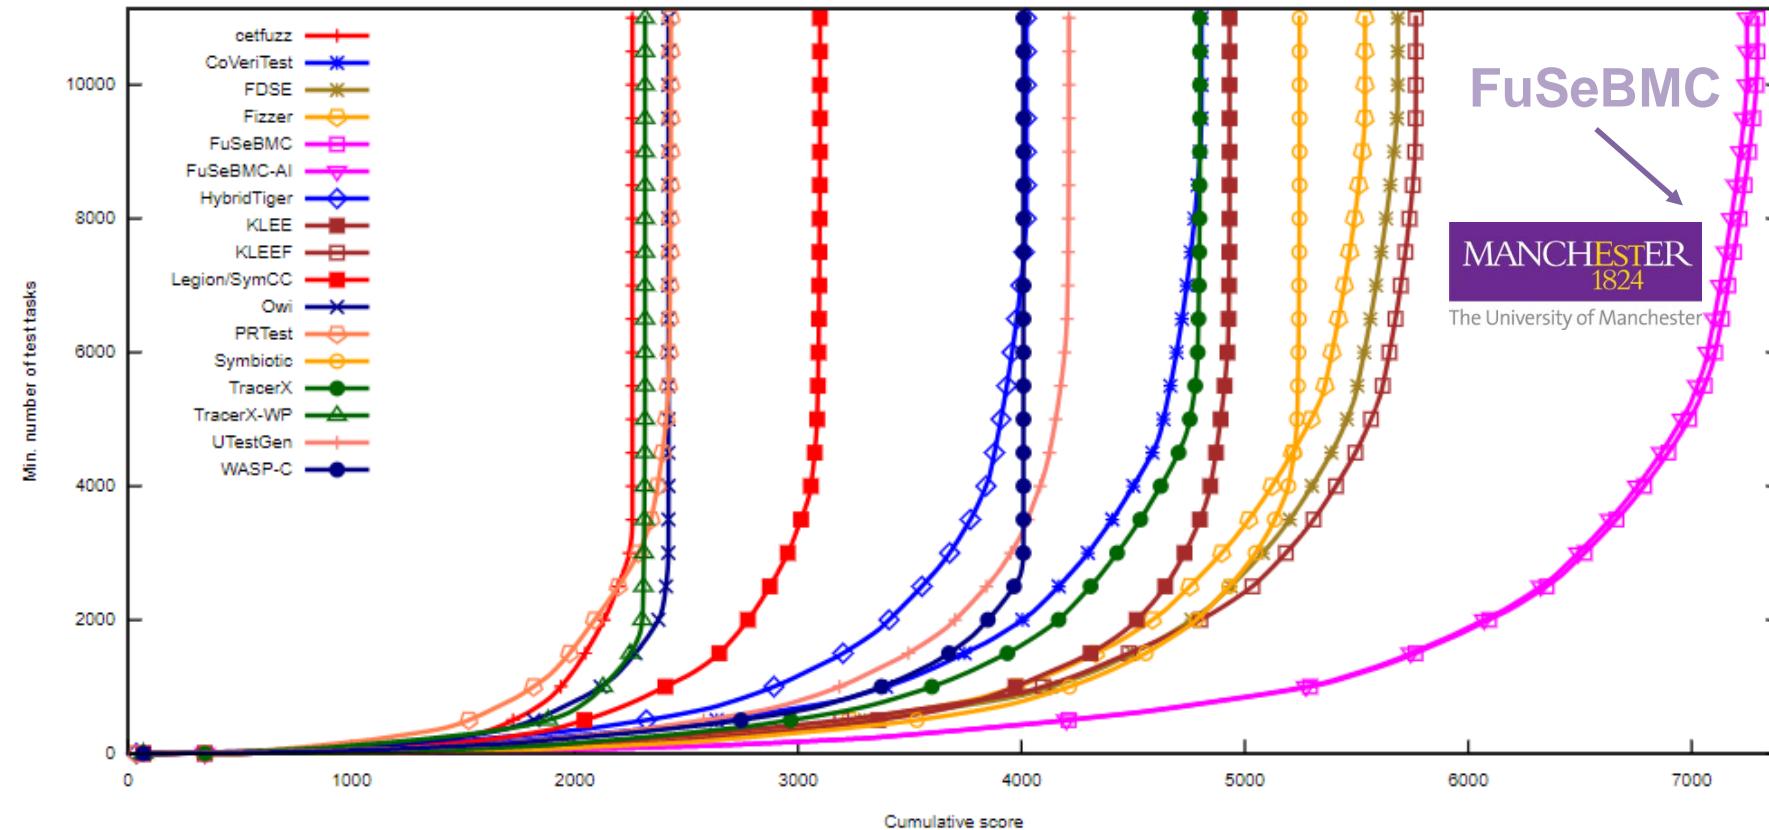
- Use **Clang** tooling infrastructure
- Employ three engines in its **reachability analysis**: **one BMC and two fuzzing engines**
- Use a **tracer** to coordinate the various engines



Agenda

- Automated program verification and testing to ensure conformance to specifications
- Advancing vulnerability detection to enhance software security beyond the state-of-the-art
- Automated program repair to improve the detection, comprehension, and fixing of software defects
- Summary and vision for future exploration

Competition on Software Testing 2024: Results of the Overall Category

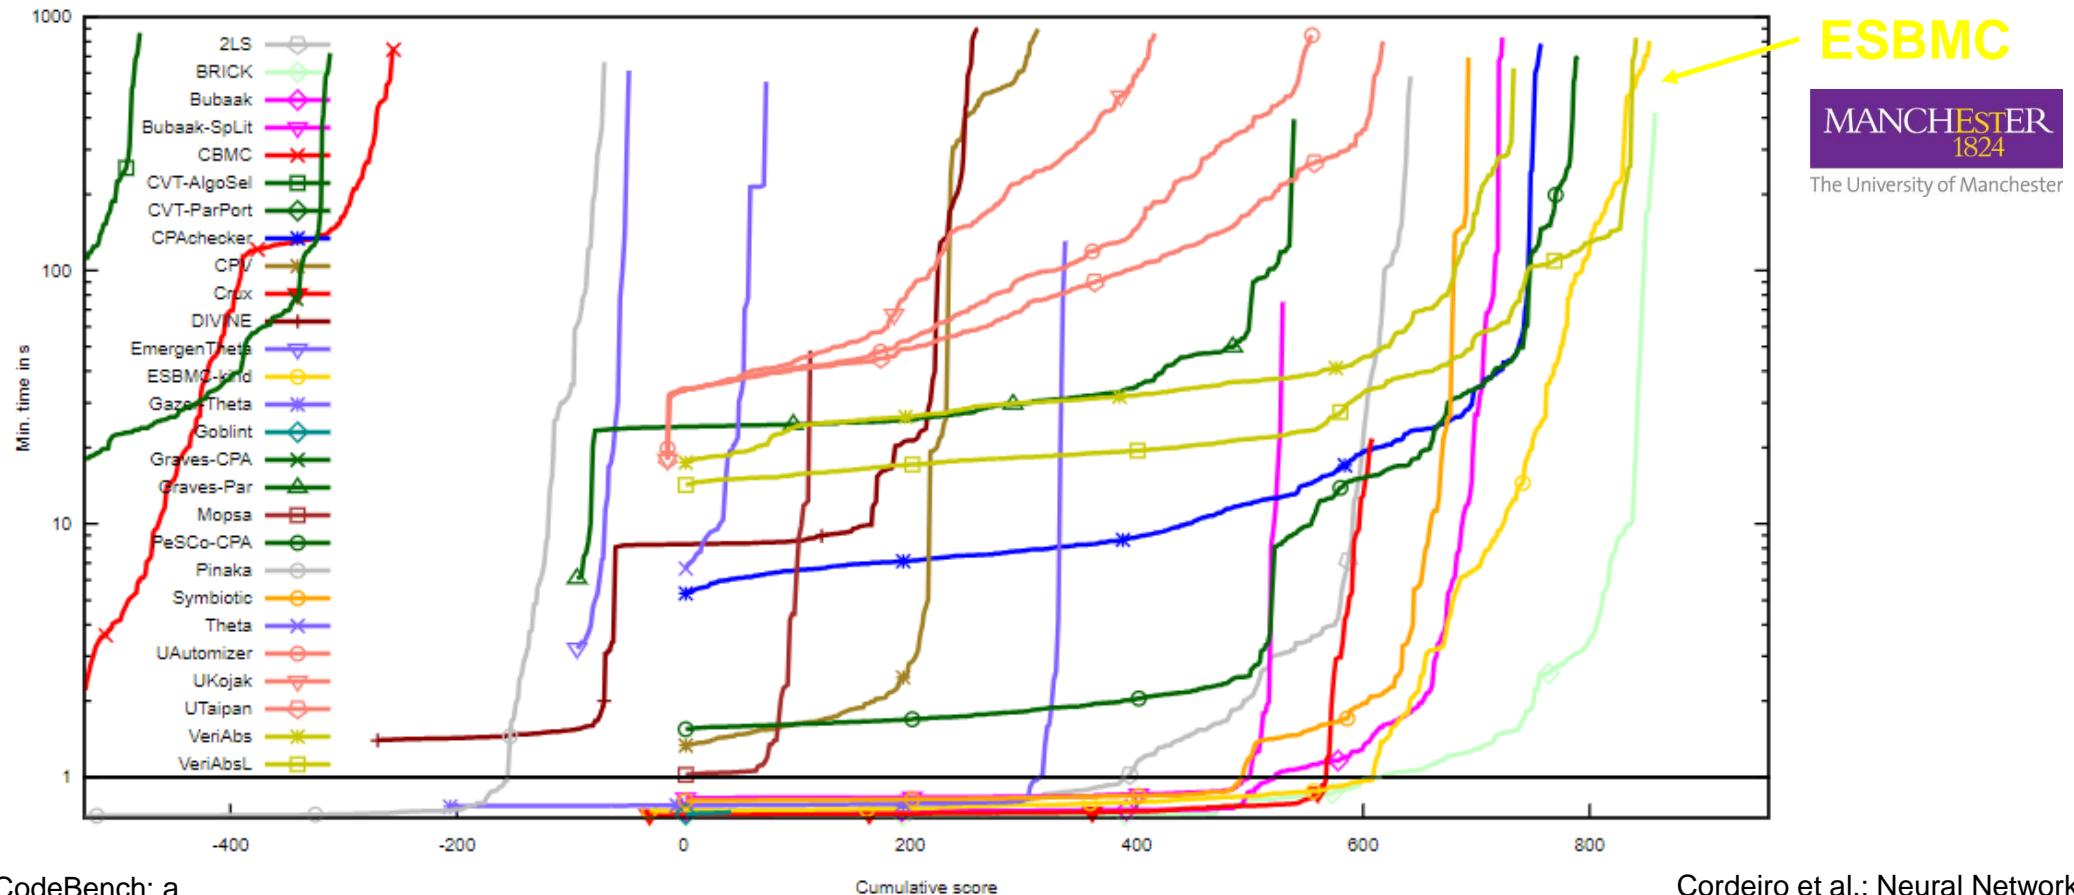


FuSeBMC achieved 3 awards: 1st place in Cover-Error, 1st place in Cover-Branches, and 1st place in Overall

From Floating-Point Programs to Neural Network Implementations



- Known ground truth, width (1-1024 neurons), depth (1-4 layers), feedforward & recurrent, 8 activation functions

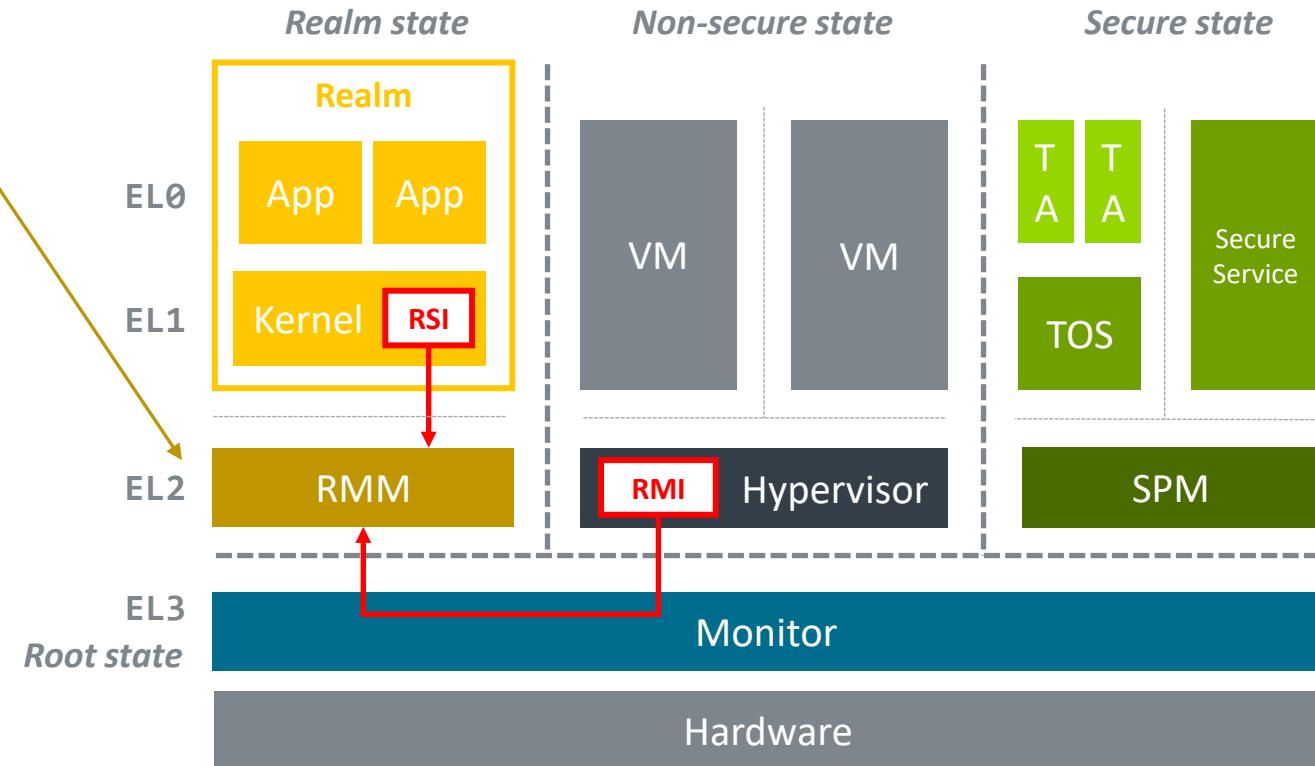


Verifying Components of Arm® Confidential Computing Architecture with ESBMC



Realm Management Monitor (RMM)

- + Provides services to Host and Realm
 - Contains no policy
 - Performs no dynamic memory allocation
- + Realm Management Interface (RMI)
 - Secure Monitor Call Calling Convention (SMCCC) interface called by Host
 - Create/destroy Realms
 - Manage Realm memory, manipulating stage 2 translation tables
 - Context switch between Realm VCPUs
- + Realm Services Interface (RSI)
 - SMCCC interface called by Realm
 - Measurement and attestation
 - Handshakes involved in some memory management flows



Arm CCA is an architecture that provides Protected Execution Environments called Realms

Verifying Components of Arm® Confidential Computing Architecture with ESBMC

- + The specification document¹ is in the style of:

- rules-based writing

R_{TMGSL}

When the state of a Granule has transitioned from P to DELEGATED and then to any other state, any content associated with P has been *wiped*.

- pre/post-condition pairs.

D3.2.5 RMI_GRANULE_DELEGATE

Delegates a Granule.

D3.2.5.1 Interface

D3.2.5.1.2 Input Values

Name	Register	Field	Type	Description
fid	X0	[63:0]	UInt64	Command FID
addr	X1	[63:0]	Address	PA of the target Granule

D3.2.5.1.3 Output Values

Name	Register	Field	Type	Description
result	X0	[63:0]	ReturnCode	Command return status

D3.2.5.2 Failure conditions

ID	Condition
gran_align	pre: !AddrIsGranuleAligned(addr) post: ResultEqual(result,RMI_ERROR_INPUT)

(— continued in the right column)

(— from the left column)

gran_bound	pre: !PaIsDelegable(addr)
	post: ResultEqual(result,RMI_ERROR_INPUT)
gran_state	pre: Granule(addr).state != UNDELEGATED
	post: ResultEqual(result,RMI_ERROR_INPUT)
gran_pas	pre: Granule(addr).pas != NS
	post: ResultEqual(result,RMI_ERROR_INPUT)

D3.2.5.3 Success conditions

ID	Post-condition
gran_state	Granule(addr).state == DELEGATED
gran_pas	Granule(addr).pas == REALM

D3.2.5.4 Footprint

ID	Value
gran_state	Granule(addr).state
gran_pas	Granule(addr).pas

- + The document is generated from a **machine-readable specification** (MRS)

¹ <https://developer.arm.com/documentation/den0137/latest>, the examples in this slide are taken when the paper was drafted.

Verifying Components of Arm® Confidential Computing Architecture with ESBMC

Test_benchmarks	esbmc multi	cbmc multi
RMI_REC_DESTROY	20	20
RMI_GRANULE_DELEGATE	safe	safe
RMI_GRANULE_UNDELEGATE	1	1
RMI_REALM_ACTIVATE	3	safe
RMI_REALM_DESTROY	15	1
RMI_REC_AUX_COUNT	1	1
RMI_FEATURES	safe	safe
RMI_DATA_DESTROY	>=24	22

```
#include <assert.h>
extern int nondet_int();
int main() {
    int m = nondet_int();
    int *n = &m;
    if((unsigned long)n >= (unsigned long)(-4095))
        assert((unsigned int)(-1 * (long)n) < 6);
    int a = -2048;
    if((unsigned long)a >= (unsigned long)(-4095))
        assert((unsigned int)(-1 * (long)a) < 6);
}
```



tautschnig commented on Jan 16

In C, pointer-to-integer conversion is implementation-defined behaviour. That should give CBMC the freedom to choose an implementation where the condition `(unsigned long)n >= (unsigned long)(-4095)` never evaluates to true.

It is, however, also right to argue that CBMC should seek to model all possible implementations. The pointer-to-integer conversion in CBMC does not currently fulfil this expectation, but we will hopefully fix this in future.



<https://github.com/diffblue/cbmc/issues/8161>

Intel Core Power Management Firmware



Intel routinely employs ESBMC to
automate firmware analysis

ESBMC has been applied to the
Authenticated Code Module, where
it found over 30 vulnerabilities

ESBMC is part of the CI pipeline for
developing microcode for the Core
family of processors

P6 Microcode Can Be Patched

Intel Discloses Details of Download Mechanism for Fixing CPU Bugs

“Taking an unusual approach to fixing bugs, Intel has implemented a microcode patch capability in its P6 processors, including Pentium Pro and Pentium II. This capability allows the microcode to be altered after the processor is fabricated, repairing bugs that are found after the processor is designed. Intel has already used this feature several times to correct minor bugs, and in the future, it may save the company from recalling CPUs if a major problem is discovered.”



WolfMQTT Verification

- **wolfMQTT library** is a client implementation of the MQTT protocol written in C for IoT devices

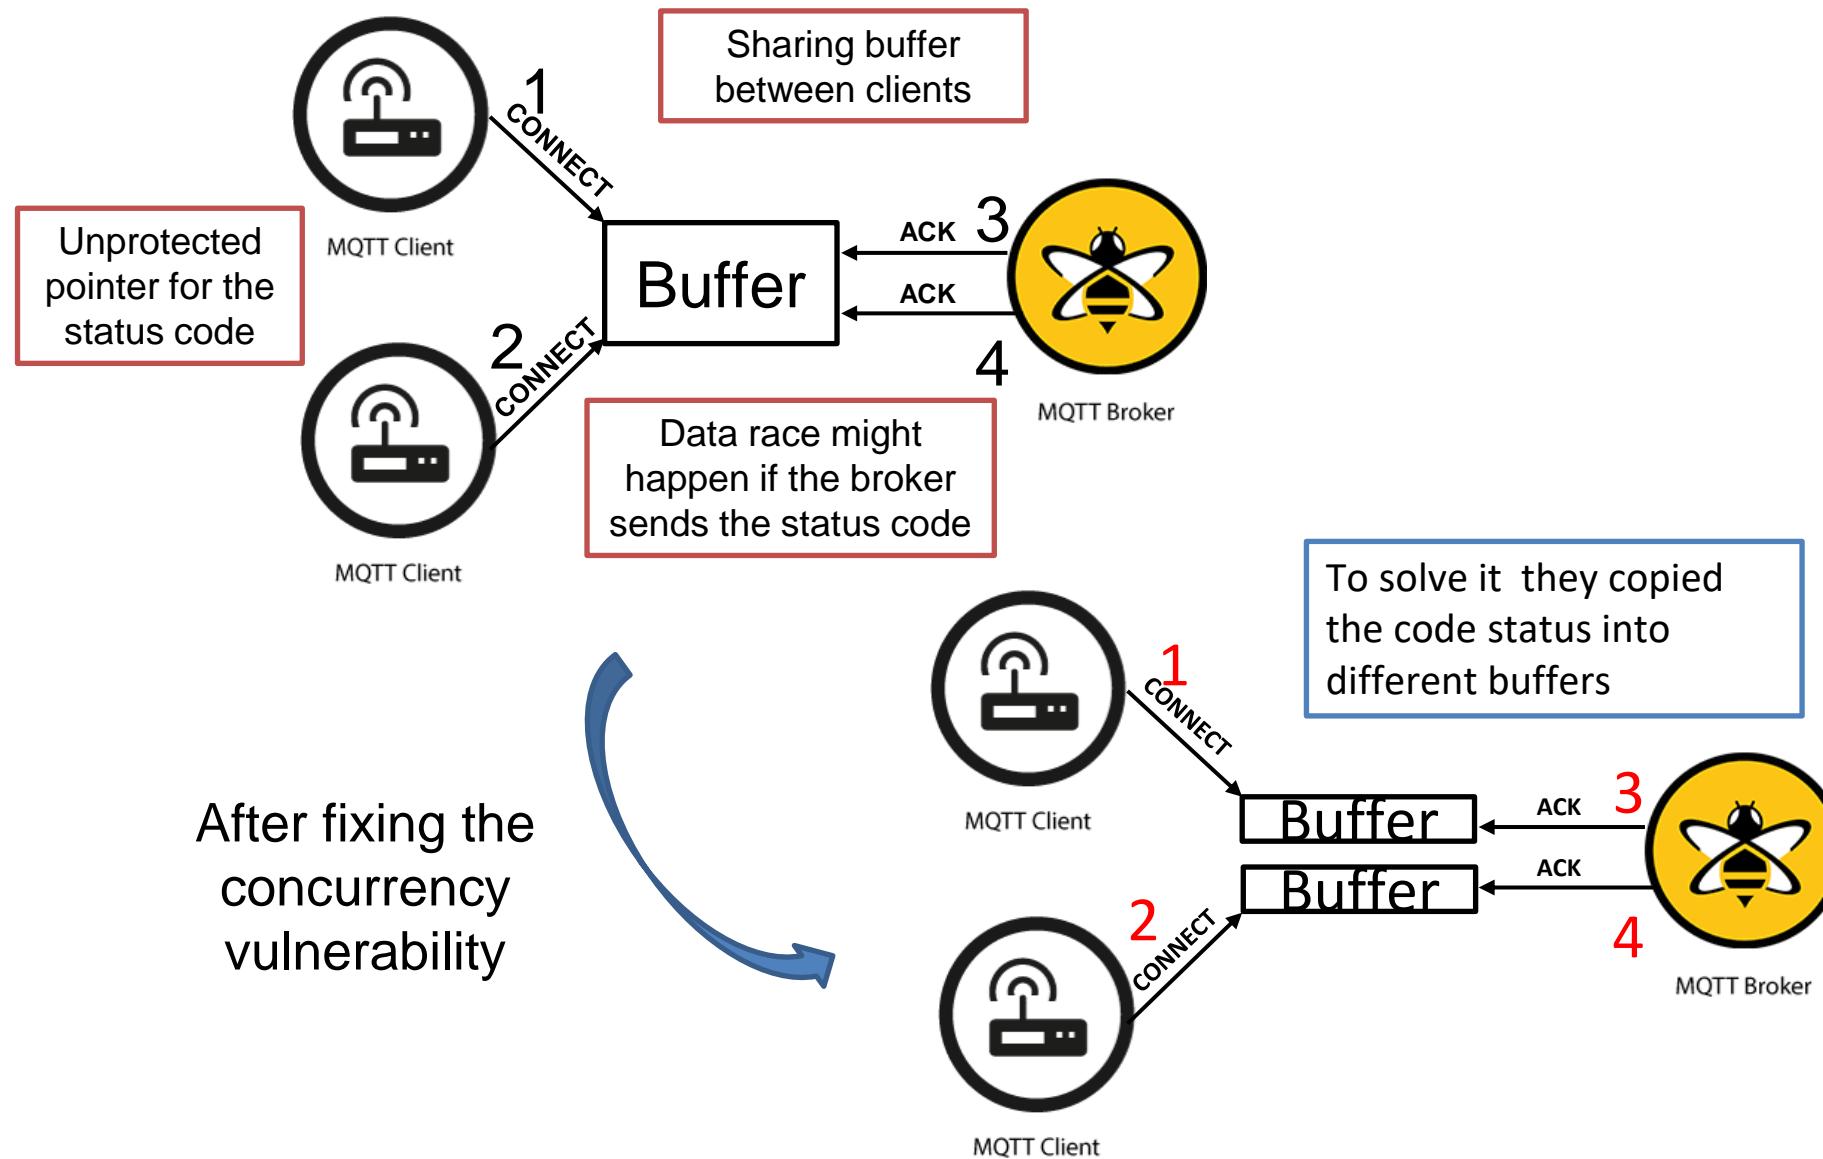
subscribe_task
and waitMessage_task are called through different threads accessing packet_ret, causing a data race in MqttClient_WaitType

Here is where the data race might happen! Unprotected pointer

```
Int main(){
    Pthread_t th1, th2;
    static MQTTCtx mqttCtx;
    pthread_create(&th1, subscribe_task, &mqttCtx);
    pthread_create(&th2, waitMessage_task, &mqttCtx)};

    static void *subscribe_task(void *client){
    .....
        MqttClient_WaitType(client, msg, MQTT_PACKET_TYPE_ANY,
        0, timeout_ms);
    .....
    static void *waitMessage_task(void *client){
    ...
        MqttClient_WaitType(client, msg, MQTT_PACKET_TYPE_ANY,
        0, timeout_ms);
    .....
    static int MqttClient_WaitType(MqttClient *client,
        void *packet_obj,
        byte wait_type, word16 wait_packet_id, int timeout_ms)
    {
    .....
        rc = wm_SemLock(&client->lockClient);
        if (rc == 0) {
            if (MqttClient_RespList_Find(client,
                (MqttPacketType)wait_type,
                wait_packet_id, &pendResp)) {
                if (pendResp->packetDone) {
                    rc = pendResp->packet_ret;
                }
            }
        }
    }
}
```

WolfMQTT Verification



Bug Report

Fixes for multi-threading issues #209

Merged embhorn merged 1 commit into wolfSSL:master from dgarske:mt_suback on 3 Jun 2021

Conversation 2 Commits 1 Checks 0 Files changed 4 +74 -48

dgarske commented on 2 Jun 2021

1. The client lock is needed earlier to protect the "reset the packet state".
2. The subscribe ack was using an unprotected pointer to response code list. Now it makes a copy of those codes.
3. Add protection to multi-thread example "stop" variable.
Thanks to Fatimah Aljaafari (@fatimahkj) for the report.
ZD 12379 and PR Data race at function MqttClient_WaitType #198

Fixes for three multi-thread issues: ... 78370ed

dgarske requested a review from embhorn 15 months ago

dgarske assigned embhorn on 2 Jun 2021

embhorn approved these changes on 3 Jun 2021

View changes

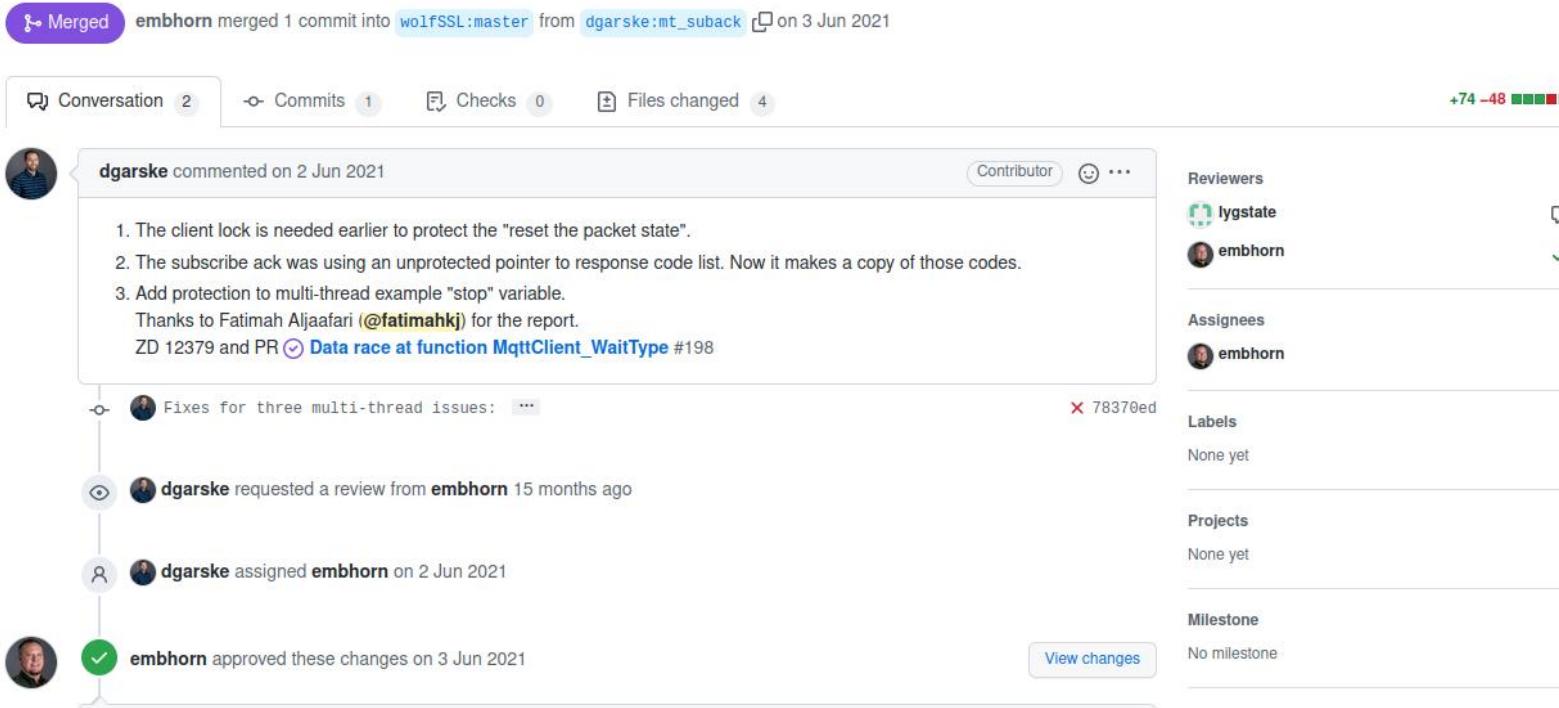
Reviewers: lygstate, embhorn

Assignees: embhorn

Labels: None yet

Projects: None yet

Milestone: No milestone



<https://github.com/wolfSSL/wolfMQTT>

Ethereum Consensus Specifications



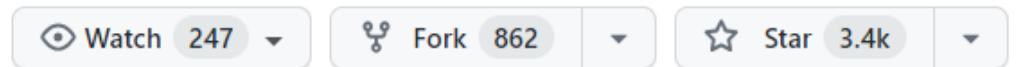
- Consensus protocol dictates how the participants in Ethereum agree on the validity of transactions and the system's state
- Git repository with **Markdown** documents describing specifications
- Infrastructure to generate **Python** libraries from Markdown

Ethereum Proof-of-Stake Consensus Specifications

chat on discord

To learn more about proof-of-stake and sharding, see the [PoS documentation](#), [sharding documentation](#) and the [research compendium](#).

This repository hosts the current Ethereum proof-of-stake specifications. Discussions about design rationale and proposed changes can be brought up and discussed as issues. Solidified, agreed-upon changes to the spec can be made through pull requests.



Contributors 148



+ 134 contributors

ESBMC-Python Benchmark

Ethereum Consensus Specification

Markdown

consensus-specs / specs / phase0 / beacon-chain.md

Preview Code Blame 1939 lines (1617 loc) · 71.4 KB

Math

```
integer_squareroot

def integer_squareroot(n: uint64) -> uint64:
    """
    Return the largest integer ``x`` such that ``x**2 <= n``.
    """
    x = n
    y = (x + 1) // 2
    while y < x:
        x = y
        y = (x + n // x) // 2
    return x

xor

def xor(bytes_1: Bytes32, bytes_2: Bytes32) -> Bytes32:
    """
    Return the exclusive-or of two 32-byte strings.
    """
    return Bytes32(a ^ b for a, b in zip(bytes_1, bytes_2))
```

eth2spec Python Library

mainnet.py ✘

```
def integer_squareroot(n: uint64) -> uint64:
    """
    Return the largest integer ``x`` such that ``x**2 <= n``.
    """
    x = n
    y = (x + 1) // 2
    while y < x:
        x = y
        y = (x + n // x) // 2
    return x

def xor(bytes_1: Bytes32, bytes_2: Bytes32) -> Bytes32:
    """
    Return the exclusive-or of two 32-byte strings.
    """
    return Bytes32(a ^ b for a, b in zip(bytes_1, bytes_2))

def bytes_to_uint64(data: bytes) -> uint64:
    """
    Return the integer deserialization of ``data`` interpreted as ``ENDIANNESS``-endian.
    """
    return uint64(int.from_bytes(data, ENDIANNESS))
```

Python Application

integer_squareroot.py ✘

```
from eth2spec.bellatrix import mainnet as spec
from eth2spec.utils.ssz.ssz_typing import (uint64)

x = uint64(16)
assert spec.integer_squareroot(x) == 4

x = uint64(25)
assert spec.integer_squareroot(x) == 5
```

ESBMC

Verification Output



Handle `integer_sqreroot` bound case #3600

Merged

hwwhww merged 3 commits into `dev` from `integer_sqreroot` 2 weeks ago

Conversation 4

Commits 3

Checks 15

Files changed 5



hwwhww commented 2 weeks ago • edited

Contributor ...

Credits to the University of Manchester Bounded Model Checking (BMC) project team: Bruno Farias, Youcheng Sun, and Lucas C. Cordeiro for reporting this issue! 🙏 100

This team is an [Ethereum Foundation ESP](#) "Bounded Model Checking for Verifying and Testing Ethereum Consensus Specifications (FY22-0751)" project grantee. They used [ESBMC model checker](#) to find this issue.

Description

`integer_sqreroot` raises `ValueError` exception when `n` is maxint of `uint64`, i.e., `2**64 - 1`.

However, we only use `integer_sqreroot` in

1. `integer_sqreroot(total_balance)`
2. `integer_sqreroot(SLOTS_PER_EPOCH)`

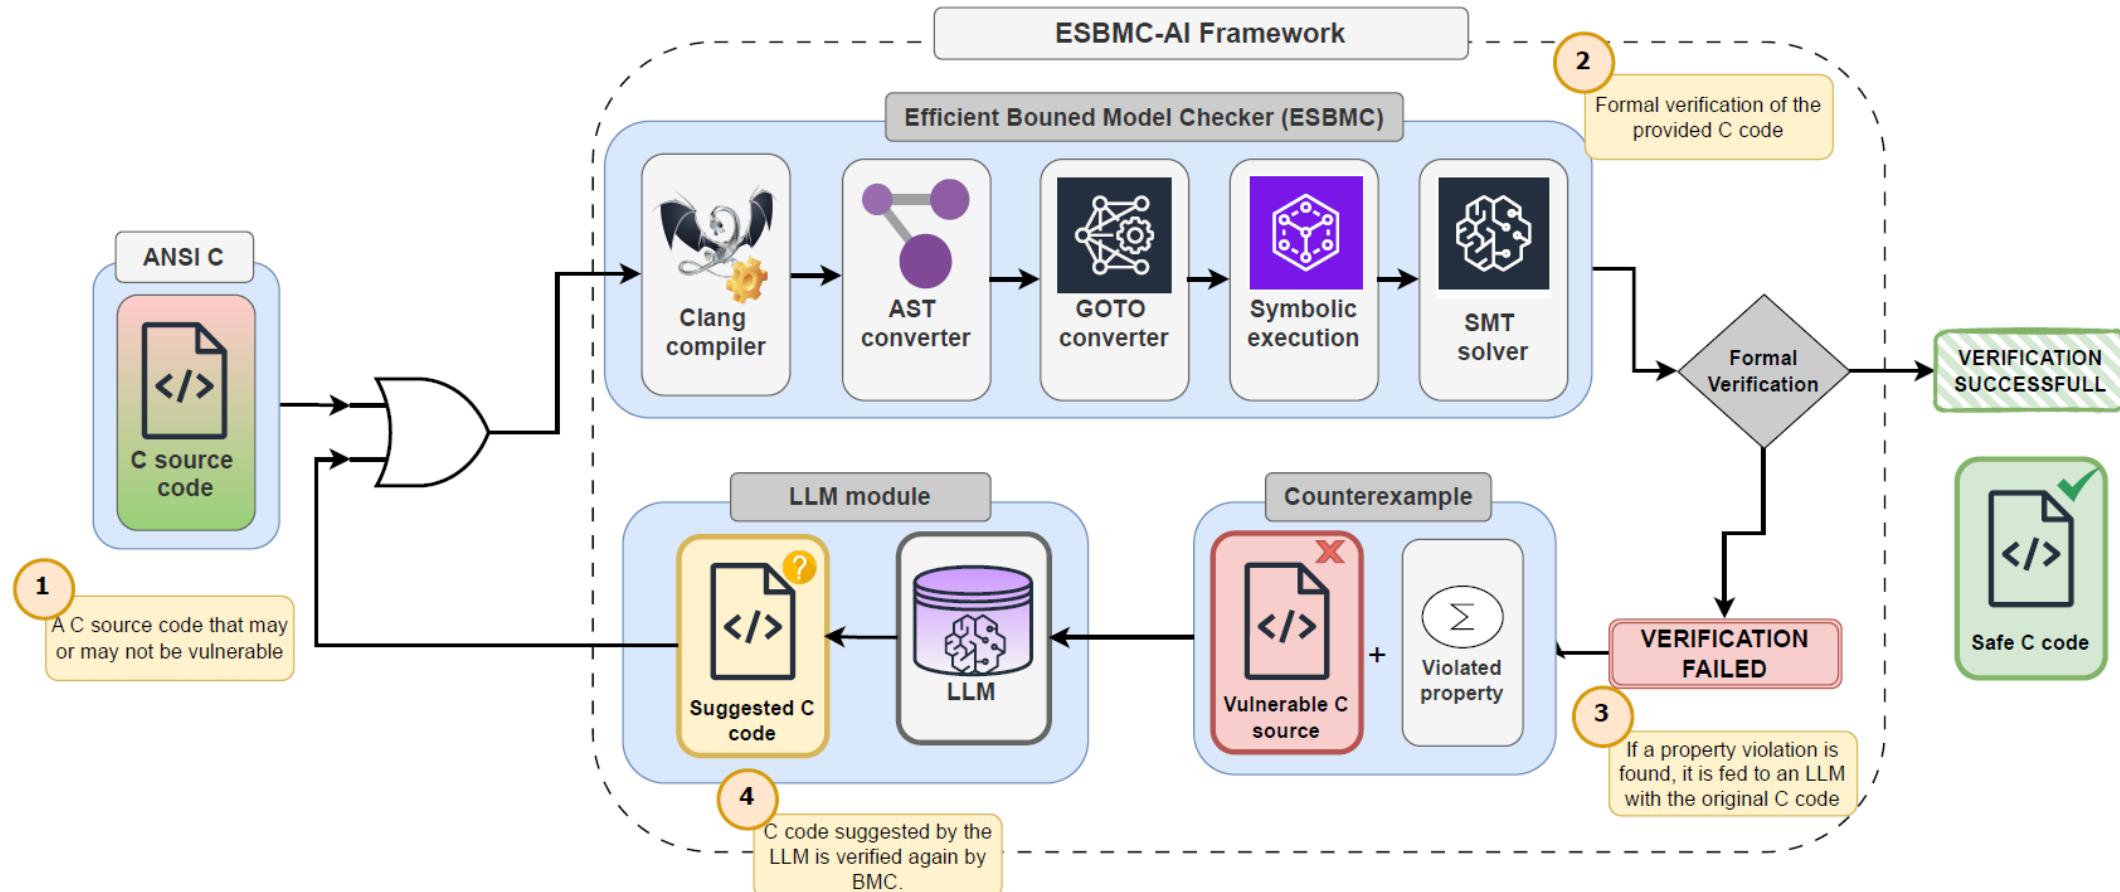
With the current Ether total supply + EIP-1559, it's unlikely to hit the overflow bound in a very long time. (↗️🔊)

That said, it should be fixed to return the expected value.

Agenda

- Automated program verification and testing to ensure conformance to specifications
- Advancing vulnerability detection to enhance software security beyond the state-of-the-art
- Automated program repair to improve the detection, comprehension, and fixing of software defects
- Summary and vision for future exploration

ESBMC-AI: Automated Program Repair



How secure is AI-generated Code: A Large-Scale Comparison of Large Language Models

Norbert Tihanyi^{1,2}, Tamas Bisztray³, Mohamed Amine Ferrag⁴, Ridhi Jain², Lucas C. Cordeiro^{5,6}

¹ Eötvös Loránd University (ELTE), Budapest, Hungary.

² Technology Innovation Institute (TII), Abu Dhabi, UAE.

³ University of Oslo, Oslo, Norway.

⁴Guelma University, Guelma, Algeria.

⁵The University of Manchester, Manchester, UK.

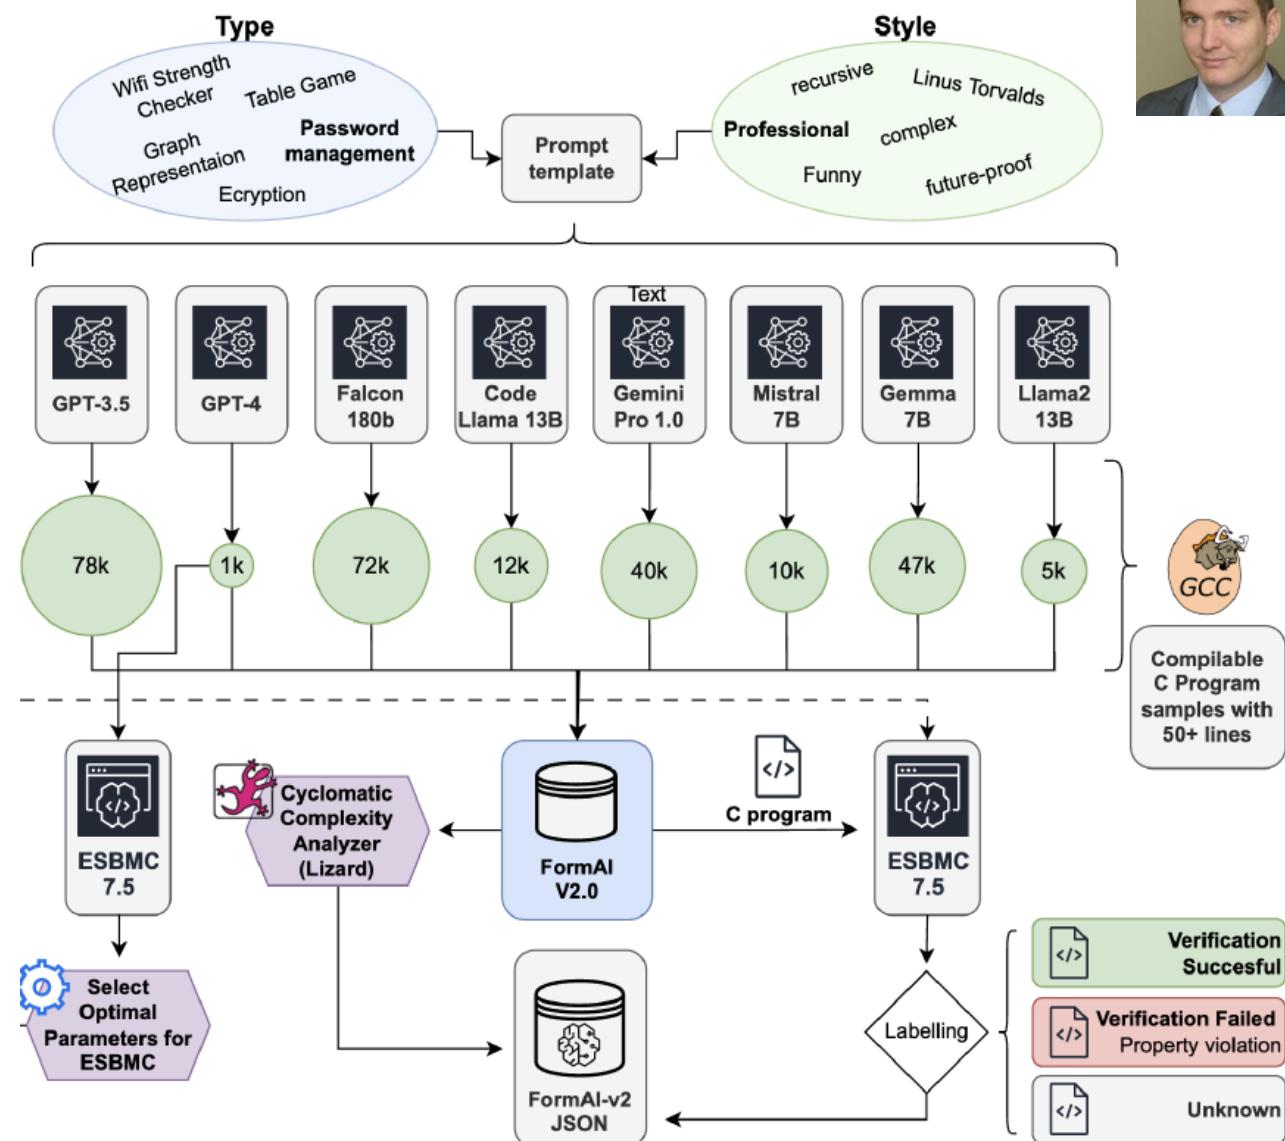
⁶Federal University of Amazonas, Manaus, Brazil.

Specs \ Datasets	Big-Vul	Draper	SARD	Juliet	Devgan	REVEAL	Diverse Vul	FormAI	FormAI v2
Language	C/C++	C/C++	Multi	Multi	C	C/C++	C/C++	C	C
Source	RW	Syn + RW	Syn + RW	Syn	RW	RW	RW	AI	AI
Dataset size	189k	1,274k	101k	106k	28k	23k	379k	112k	150k
Vul. Snippets	100%	5.62%	100%	100%	46.05%	9.85%	7.02%	51.24%	61%
Multi. Vulns.	✗	✓	✗	✗	✗	✗	✗	✓	✓
Compilable	✗	✗	✓	✓	✗	✗	✗	✓	✓
Granularity	Func	Func	Prog	Prog	Func	Func	Func	Prog	Prog
Class. Type	CVE CWE	CWE	CWE	CWE	CVE	CWE	CWE	CWE	CWE
Avg. LOC.	30	29	114	125	112	32	44	79	82
Labelling Method	P	S	B/S/M	B	M	P	P	F	F

Legend:

Multi: Multi-Language Dataset, RW: Real World, Syn: Synthetic, AI: AI-generated,
 Func: Function level granularity, Prog: Program level granularity,
 CVE: Common Vulnerabilities and Exposures, CWE: Common Weakness Enumeration,
 P: GitHub Commits Patching a Vulnerability, S: Static Analyzer,
 B: By Design Vulnerable, F: Formal Verification with ESBMC, M: Manual Labeling

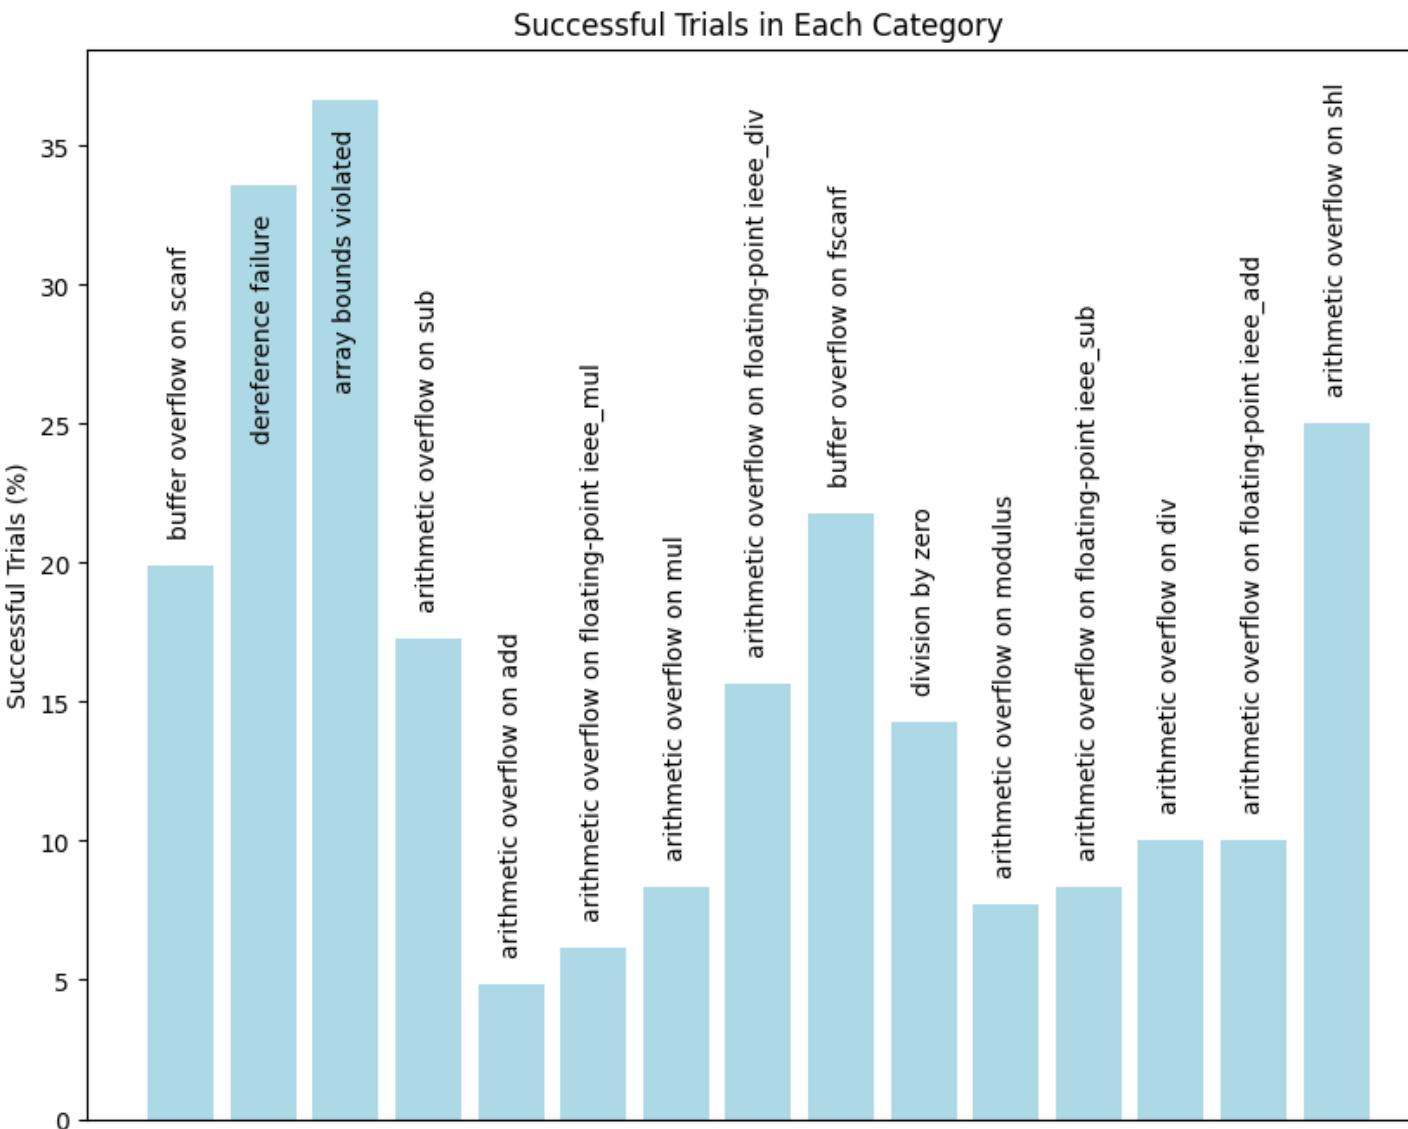
<https://doi.org/10.1007/s10664-024-10590-1>



Network Management, Table Games, Wi-Fi Signal Strength Analyzer, QR code reader, Image Steganography, Pixel Art Generator, Scientific Calculator Implementation, and Encryption, string manipulation, etc.



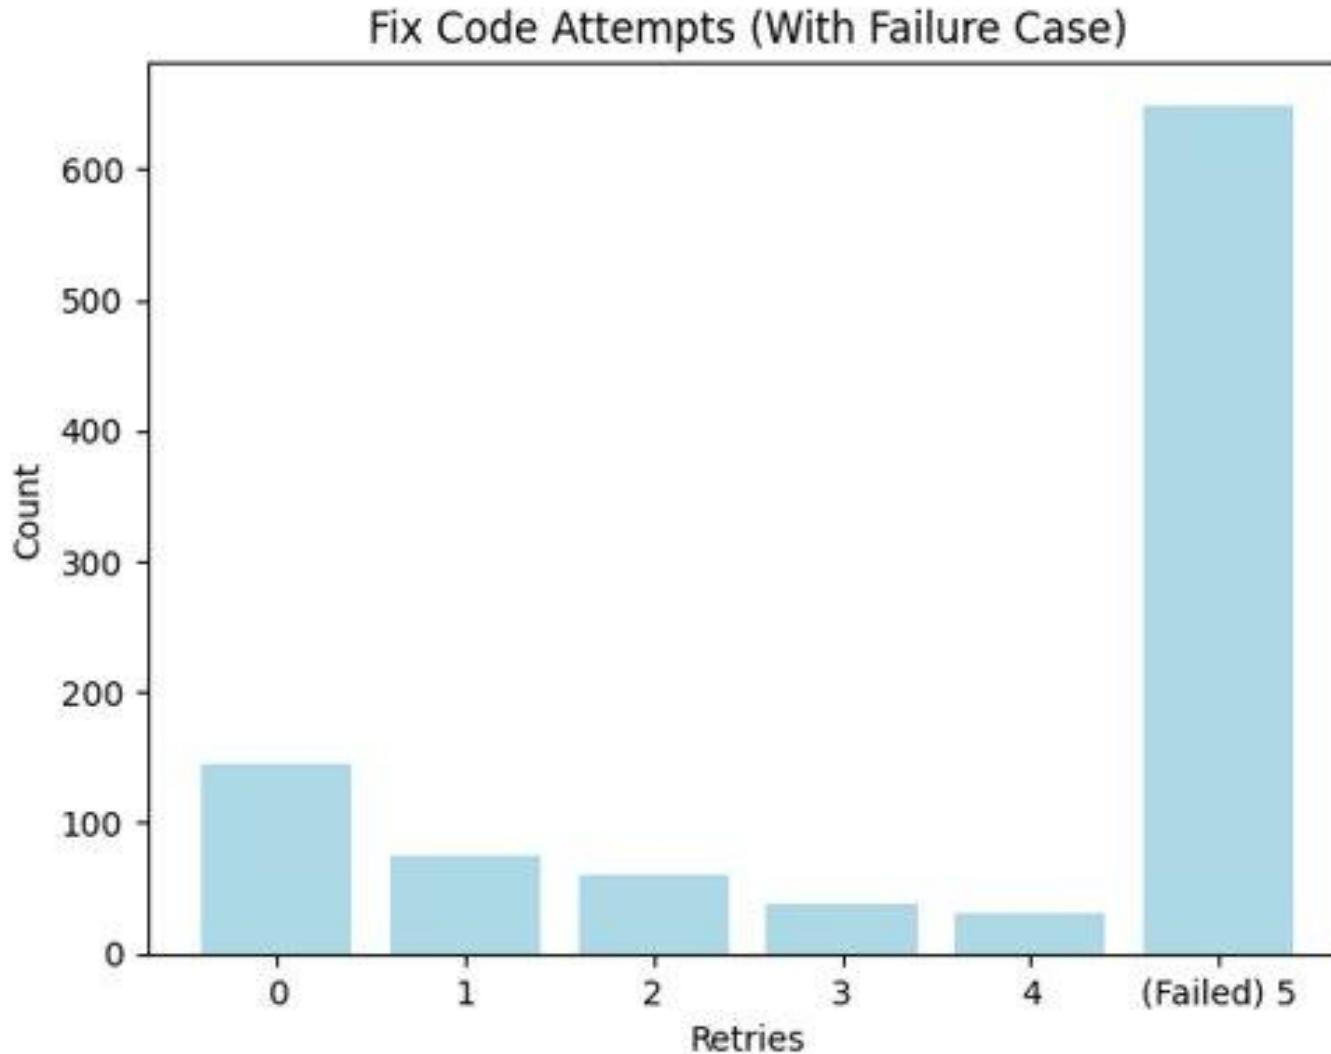
ESBMC-AI Fix Code Mode (FCM) Performance



- **Built the formAI dataset with 112k C programs**
- Randomly selected 1k vulnerable C programs
- Repaired 35.5% programs
- Lowest category was arithmetic overflow (~5%)
- Highest category was array out of bounds (~36%)
- **Generic prompts (room for improvement)**

Charalambous, Y. et al.: A New Era in Software Security: Towards Self-Healing Software via Large Language Models and Formal Verification. Under review at AST 2025.

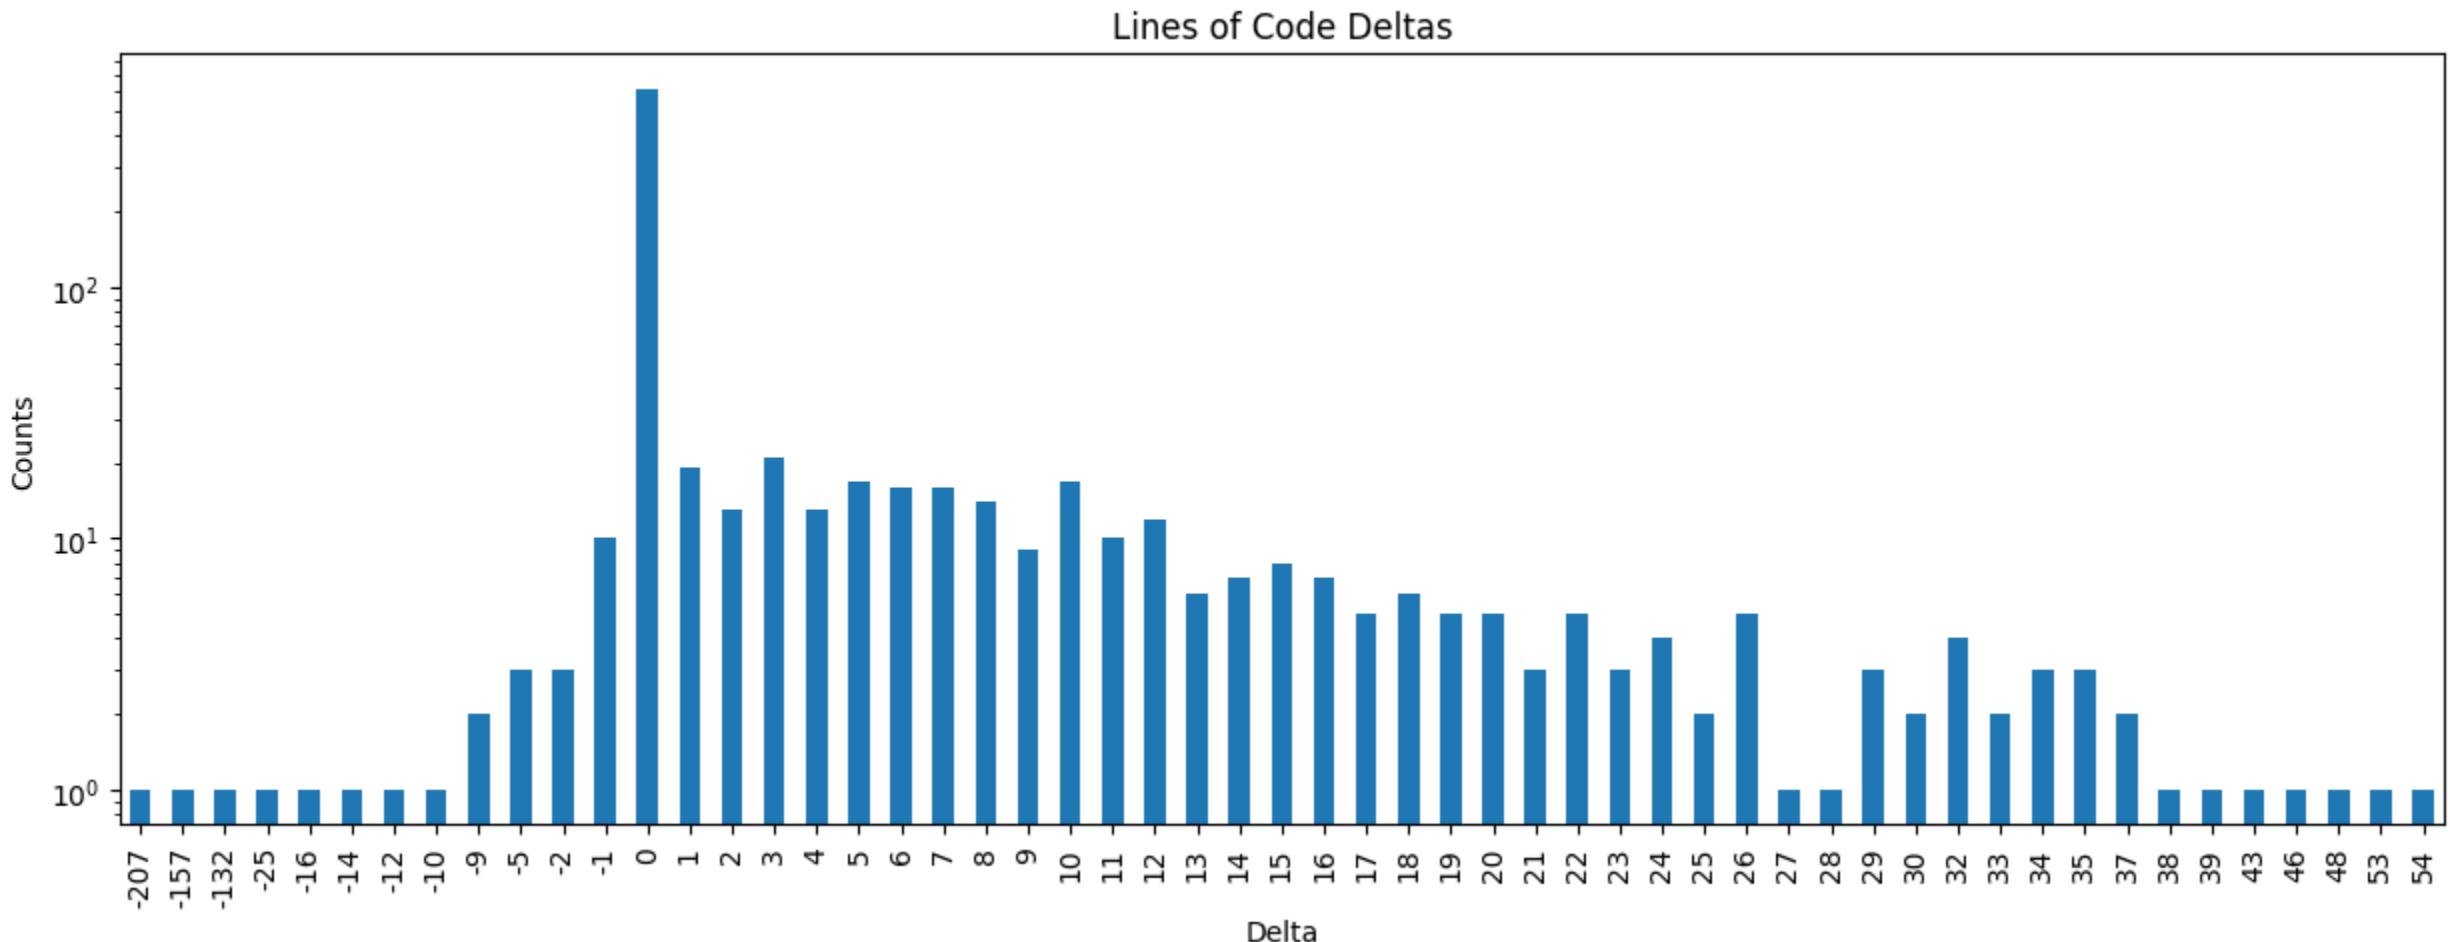
ESBMC-AI Fix Code Mode (FCM) Performance



- **Built the formAI dataset with 112k C programs**
- Randomly selected 1k vulnerable C programs
- Repaired 35.5% programs
- Lowest category was arithmetic overflow (~5%)
- Highest category was array out of bounds (~36%)
- **Generic prompts (room for improvement)**

Charalambous, Y. et al.: A New Era in Software Security: Towards Self-Healing Software via Large Language Models and Formal Verification. Under review at AST 2025.

ESBMC-AI Fix Code Mode (FCM)

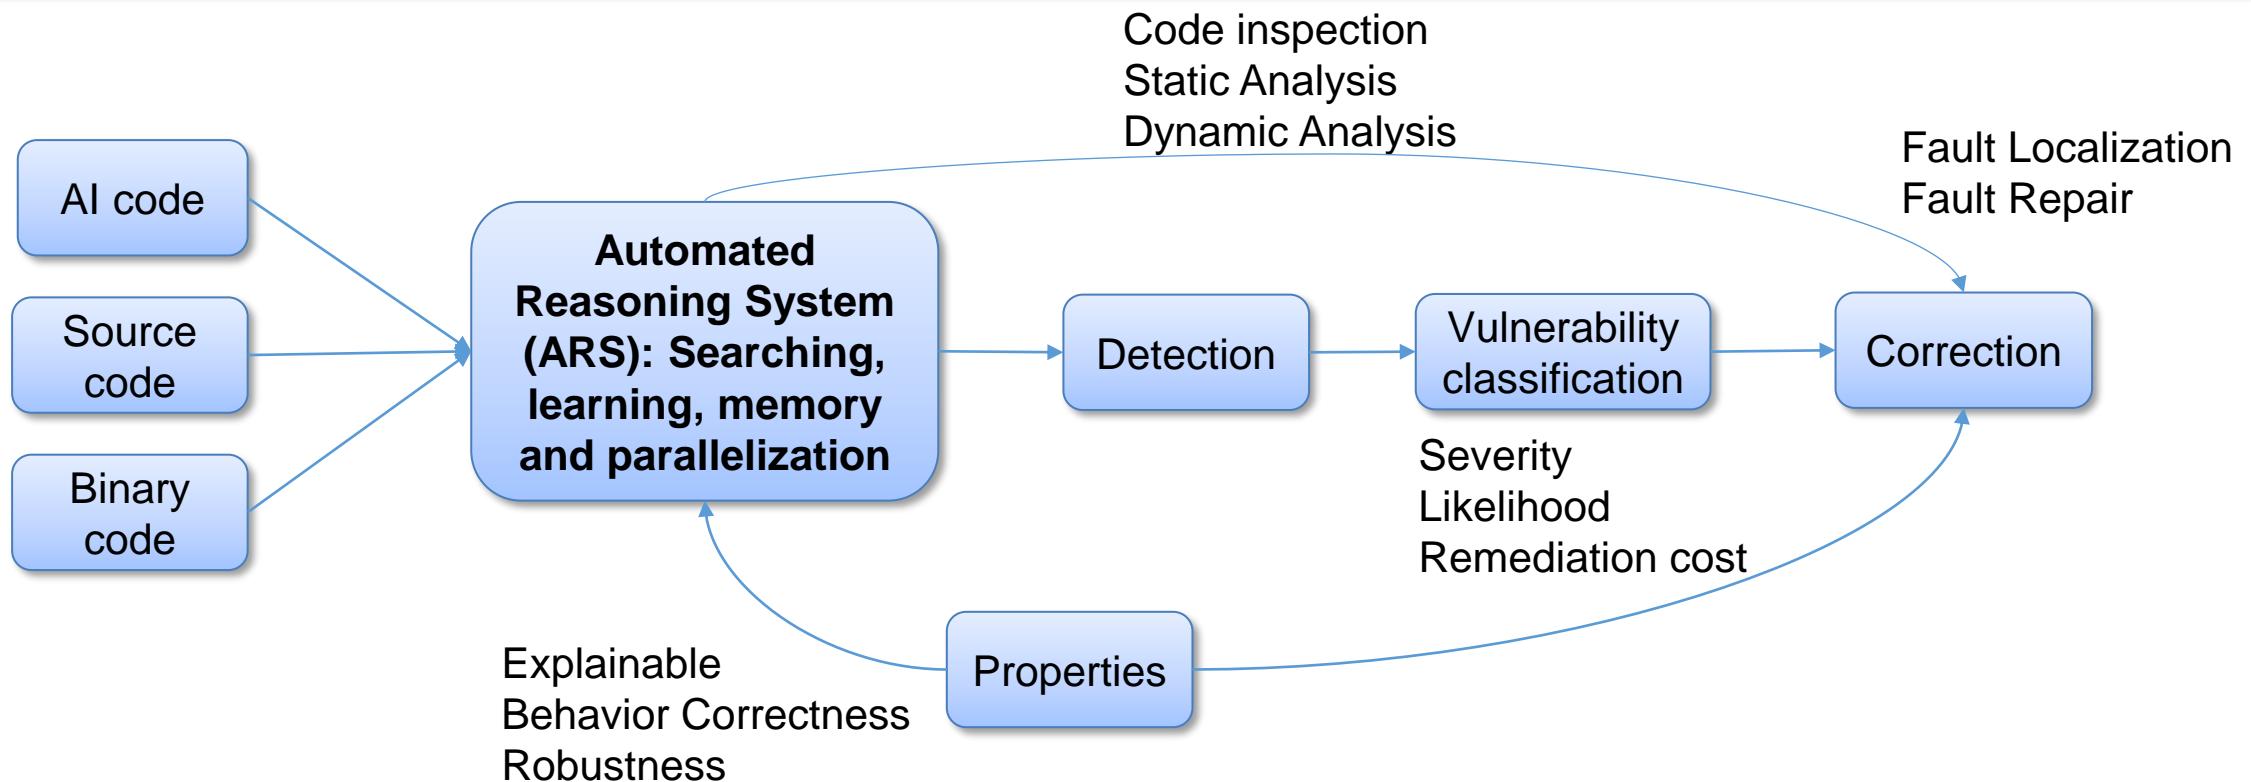


Agenda

- Automated program verification and testing to ensure conformance to specifications
- Advancing vulnerability detection to enhance software security beyond the state-of-the-art
- Automated program repair to improve the detection, comprehension, and fixing of software defects
- Summary and vision for future exploration

Vision: Building Trustworthy Software and AI Systems

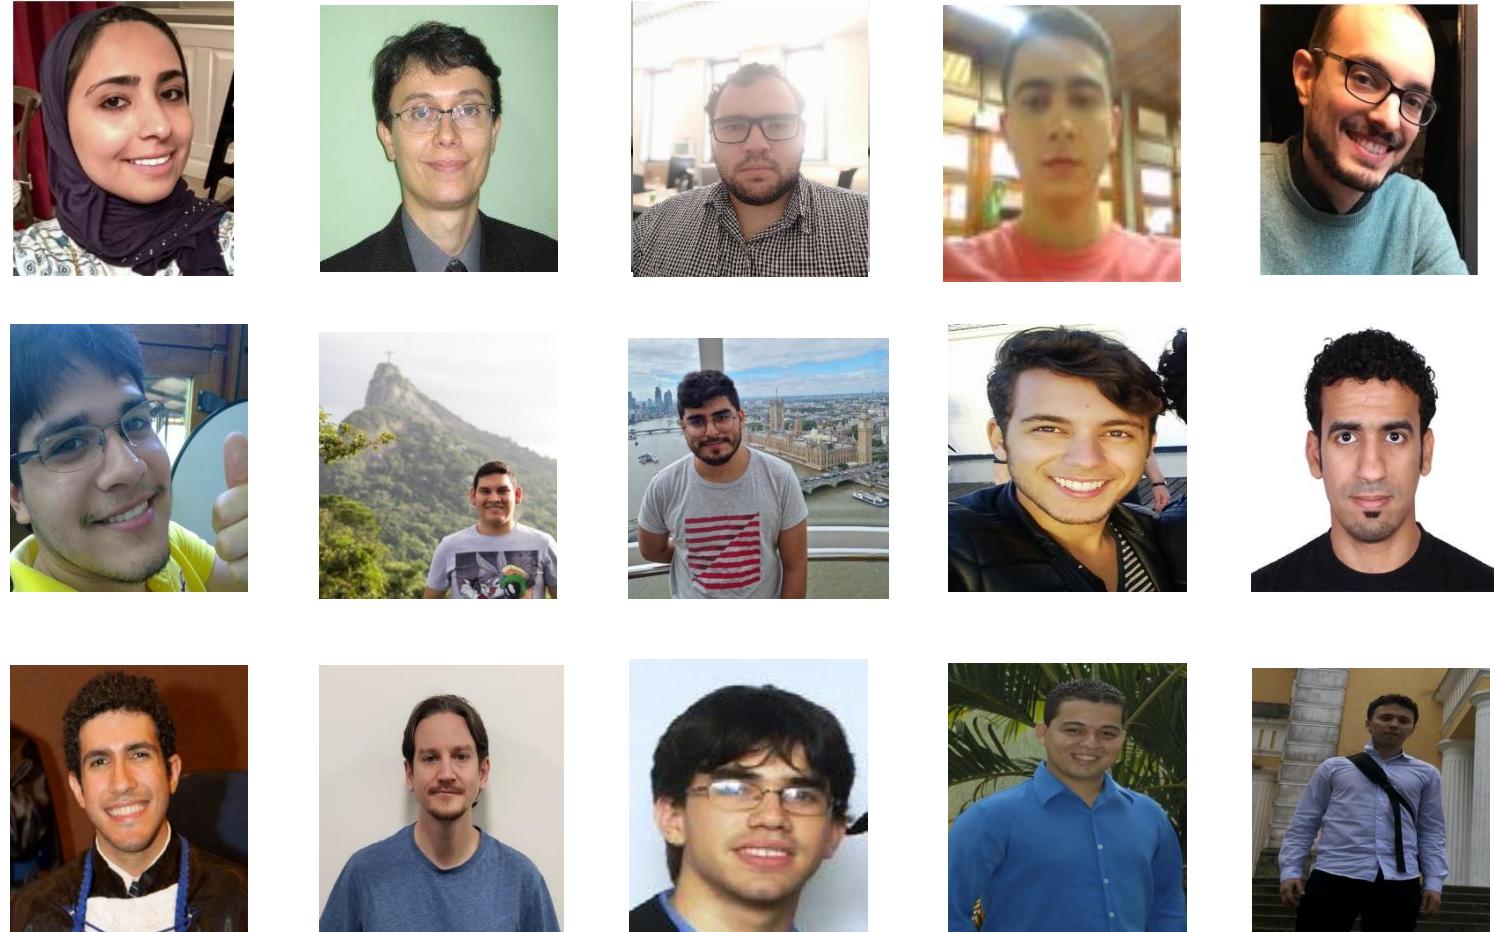
Develop an automated reasoning system for **safeguarding software and AI systems** against vulnerabilities in an increasingly digital and interconnected world



(Real) Impact: Students and Contributors

- 5 PhD theses
- 30+ MSc dissertations
- 30+ final-year projects
- GitHub:
 - 43 contributors
 - 24,268 commits
 - 315 stars
 - 101 forks

<https://github.com/esbmc/esbmc>



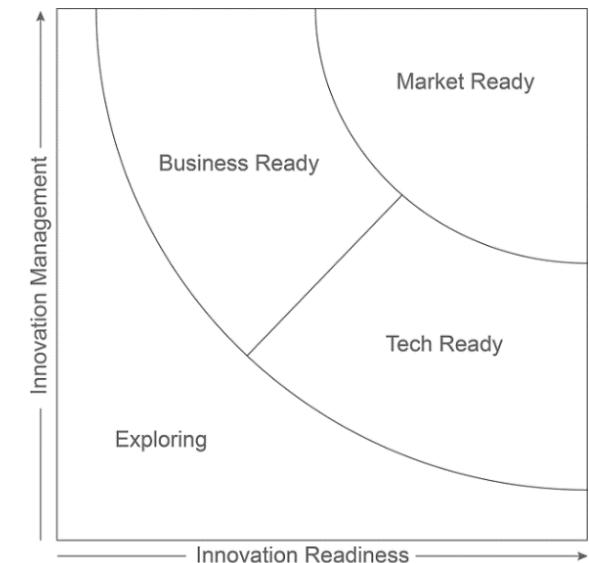
Impact: Awards and Industrial Deployment

- **Distinguished Paper Award** at ICSE'11 and ASE'24
- **Best Paper Award** at SBESC'15
- **Most Influential Paper Award** at ASE'23
- **Best Tool Paper Award** at SBSeg'23
- **46 awards** from intl. competitions on SW verification/testing at **TACAS/FASE**
 - Bug Finding and Code Coverage 
- **Intel** deploys **ESBMC** in production as one of its verification engines for **verifying firmware in C**
- **Nokia** and **ARM** have found **security vulnerabilities** in **C/C++ software**
- **Funded by the government** (EPSRC, British Council, Royal Society, CAPES, CNPq, FAPEAM) and **industry** (Intel, Motorola, Samsung, Nokia, ARM)
- **Potential spin-out** about building trustworthy software and AI systems

The European Commission recognised our code verification framework as an outstanding innovation

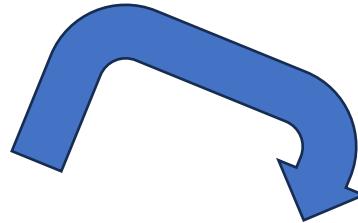


- *"We believe that your organisation's inclusion in this initiative could open up new opportunities for you to partner with business or academic organisations and trigger interest from potential customers or investors in your innovations"*
 - **Innovation Title:** ELEGANT code verification mechanisms;
 - **Market Maturity of the Innovation:** Exploring
 - **Market Creation Potential of the innovation:** High





- UKRI IAA with SES Software Escrow
- Identifying security vulnerabilities in code
- FuSeBMC has secured more than £120,000 in funding from Innovate UK and the Innovation Factory, demonstrating external support for its innovation potential

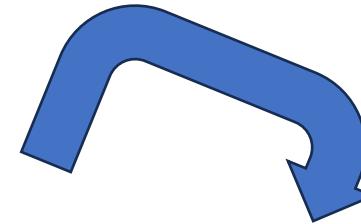


Software spin-out wins funding from Innovation Factory

Wednesday, February 14, 2024 [Research](#)



Commercial Impact

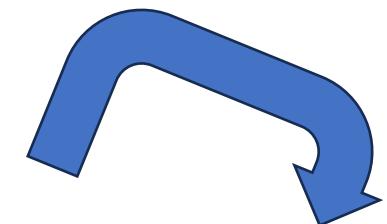


CyberASAP

The only pre-seed accelerator programme in the cybersecurity ecosystem, CyberASAP (Cyber Security Academic Startup Accelerator Programme) plays a unique and vital role in supporting cyber security innovation and commercialisation.



<https://iuk.ktn-uk.org/programme/cyberasap/>



VeriBee

<https://www.veribee.co/>

Advisory Committee by Deeptech Labs: £200,000

S3 Research Group

Security is one of the priority growth areas in the CS department's ~£10m expansion plan, and we already have **two new L/SL and one Professor posts advertised**



L. Cordeiro
(Software Security)



R. Banach
(Formal Methods)



M. Mustafa
(Applied Cryptography)



N. Zhang
(IoT Security)



E. Manino
(AI Security)



B. Magri
(Cryptography)



Y. Sun
(Trustworthy AI)



D. Dresner
(System Governance)



A. Creswell
(Cyber and AI)



G. Smith
(Former Director
General for
Technology at GCHQ)

Collaborators across the University



K. Korovin
(Automated Reasoning)



G. Reger
(Automated Reasoning)



C. Kotselidis
(Computer Architecture)



Meropi Tzanetakis
(Digital Criminology)



Kaled Alshmrany
(Software Testing)



M. Luján
(Computer Architecture)



J. Goodacre
(Computer Architecture)



P. Olivier
(Computer Architecture)



Richard Allmendinger
(Applied AI)

S3 Research Objectives

We develop state-of-the-art algorithms, methods, and protocols to address **security** and **privacy** in **networked and distributed system** environments, and tools to build verifiable, **trustworthy software and AI/ML systems**



21 March 2024

University of Manchester recognised as Academic Centre of Excellence in Cyber Security Research

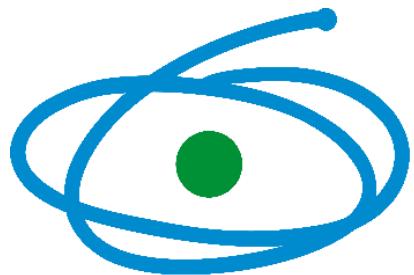
Acknowledgements



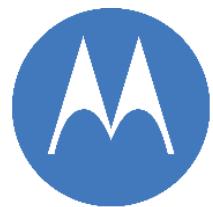
Engineering and Physical Sciences
Research Council



UK Research
and Innovation



CAPES



motorola



ethereum

FLEXTRONICS®

NOKIA

intel®

arm